

# Personal Finance Daily Expenses Tracking Prediction Using Machine Learning

## I. Introduction

This project investigates the use of machine learning for predicting daily expenses. Although no singular hypothesis was defined at the outset, the objective is to evaluate how accurately a model can forecast daily expenditures using a custom-generated dataset. The experiments focus on exploring the predictive power of engineered temporal features—such as day of the week, weekend indicators, lag features, and rolling averages—and assessing the impact of these features on model performance. In addition, unsupervised clustering methods are applied to reveal underlying patterns in spending behavior.

## II. Dataset Documentation

### A. Dataset Generation

The dataset was generated using a custom Python script (see *generate\_data.py* in the appendix) that simulates a full year (365 days) of financial transactions for a student. The simulation captures both income and various expense categories, such as:

- **Income:** Includes scholarship payments, parental support, part-time job income (e.g., from Burger King and office assistant work).
- **Expenses:** Includes fixed monthly costs (e.g., transport, entertainment subscriptions, groceries) and daily variable expenses (meals, coffee, dinner, and laundry), as well as occasional special events (e.g., weekend dining out) and gym-related spending.

### III. Data Composition and Characteristics

- **Time Frame:** January 1, 2025, to December 31, 2025.
- **Features:**

<b>Date</b>	The transaction date.
<b>Category</b>	Labels such as: Income, Transport, Entertainment, Groceries, Meal, Coffee, Food & Drink, Gym, and Laundry.
<b>Amount_NTD</b>	Transaction amounts in New Taiwan Dollars.
<b>Description, Payment_Method, Time</b>	Additional metadata.

- **Generation Process:**

A biased random function was used to simulate realistic values (e.g., amounts ending in 0 or 5 are favored). The script includes fixed monthly incomes and

recurring expenses, while daily expenditures are generated with variability based on conditions (e.g., gym days, weekends).

- **Dataset Access:**

The dataset is stored in a CSV file ([daily\\_expenses.csv](#)) and is available on my GitHub [repository](#).

## IV. Methodology

### A. Supervised Learning

To predict daily expenses, we employed a Random Forest regression model using the open-source scikit-learn library ([scikit-learn.org](#)). Our supervised learning pipeline was implemented in Python and leveraged other widely adopted libraries such as Pandas and NumPy for data manipulation, and Matplotlib and Seaborn for visualization ([pandas.pydata.org](#), [numpy.org](#), [matplotlib.org](#), [seaborn.pydata.org](#)).

- **Feature Engineering:**

Several features were derived from the raw date:

1. **DayOfWeek:** Integer representing the day of the week (Monday=0, ..., Sunday=6).
2. **IsWeekend:** Binary indicator (1 if Saturday or Sunday, else 0).
3. **Month and Day:** Extracted from the date.
4. **Lag1:** The expense value from the previous day, capturing short-term dependencies.
5. **Rolling7:** A 7-day moving average of expenses to capture trends.
6. **LogExpense:** The logarithm of daily expenses (after adding 1) to reduce the impact of outliers.

- **Model Training and Evaluation:**

The Random Forest model was trained using different feature sets:

1. **Feature Set A:** Without Lag1 and Rolling7.
  2. **Feature Set B:** With Lag1 only.
  3. **Feature Set C:** With Rolling7 only.
  4. **Feature Set D:** With both Lag1 and Rolling7.
- A grid search with cross-validation (using TimeSeriesSplit) was performed for hyperparameter tuning. Evaluation metrics included RMSE, with additional analysis of feature importance.

Additionally, our pipeline incorporated standard data resampling techniques to ensure a balanced representation of the temporal data. Dimensionality reduction, specifically Principal Component Analysis (PCA), was applied to visualize high-dimensional features in two dimensions. This aided in both exploratory data analysis and in validating the clustering results.

**B. Unsupervised Learning**

Clustering techniques were employed to explore spending behavior patterns:

- **K-Means Clustering:**  
Multiple metrics (Elbow method, Silhouette score, Davies–Bouldin Index, Calinski–Harabasz Index) were used to evaluate the optimal number of clusters. Experiments revealed variability in the optimal k, with the Elbow method suggesting  $k = 6$ .
- **Gaussian Mixture Model (GMM):**  
GMM was applied to obtain probabilistic cluster assignments and confidence scores, providing another perspective on the grouping of daily expenses.

The use of PCA for dimensionality reduction was critical in visualizing both the K-Means and GMM clustering results, as it enabled us to project the multi-dimensional feature space onto two principal components.

**V. Experiments Results & Discussion**

**A. Supervised Learning Experiments**

Features	RMSE	Cross-Validation Mean RMSE
Feature Set A (No Lag1/No Rolling7)	38.67	0.1683
Feature Set B (With Lag1 Only)	39.34	0.1655
Feature Set C (With Rolling7 Only)	37.28	0.1548
Feature Set D (With Both Lag1 & Rolling7)	36.20	0.1540

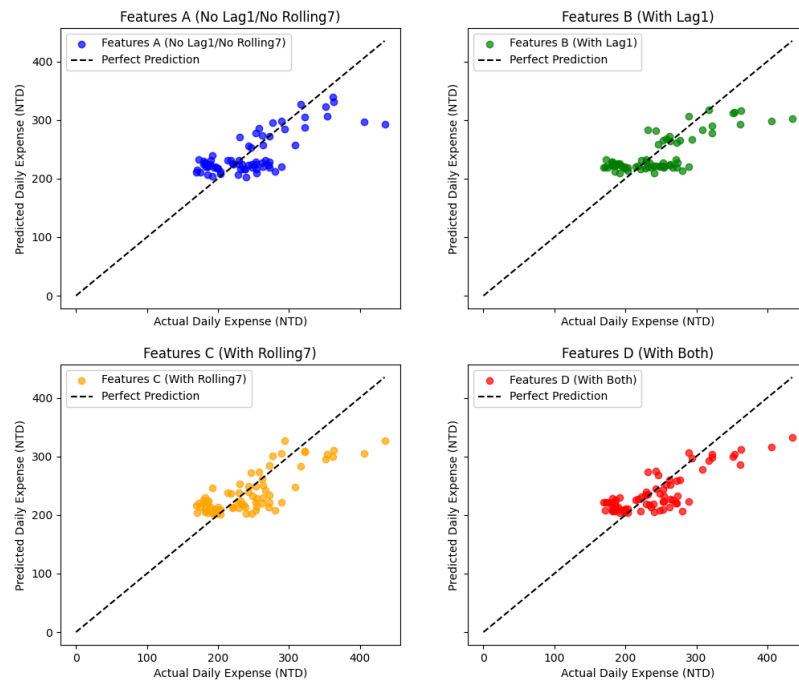
**Table 1. Random Forest Regression Performance by Different Feature Set**

In this experiments, I evaluated a Random Forest regression model using four distinct feature sets to predict daily expenses. Feature Set A, which excluded both lagged and rolling features, produced an RMSE of approximately 38.66, while incorporating a lag feature in Feature Set B resulted in a slight improvement with an RMSE near 39.34. Notably, Feature Set C—using only a 7-day rolling average—yielded a lower RMSE of about 37.28, and the combination of both lag (Lag1) and rolling average (Rolling7) in Feature **Set D** achieved the best performance with an **RMSE** of roughly **36.20**. Cross-validation confirmed the model’s stability, and feature importance analysis (Table 2) consistently highlighted **DayOfWeek** and **IsWeekend** as the most influential predictors, with the added features enhancing the prediction by capturing temporal trends.

Feature	Set A	Set B	Set C	Set D
DayOfWeek	0.406320	0.406350	0.335043	0.292754
IsWeekend	0.282965	0.314623	0.239466	0.229379
Month	0.144396	0.064919	0.078241	0.054489
Day	0.166320	0.073103	0.108969	0.079247
Lag1	-	0.141005	-	0.130432
Rolling7	-	-	0.238281	0.213698

**Table 2. Feature Importances for Random Forest with Different Feature Sets**

In addition, unsupervised learning experiments using K-Means clustering and Gaussian Mixture Models revealed that the dataset naturally segmented into distinct groups; while the Elbow method suggested an optimal k of 6, alternative metrics provided varied insights, and GMM clustering further confirmed the presence of six distinct clusters with high confidence, each exhibiting unique spending profiles. These findings underscore that careful feature engineering—particularly the integration of temporal dynamics—significantly improves predictive accuracy and offers valuable insights into underlying spending behaviors.

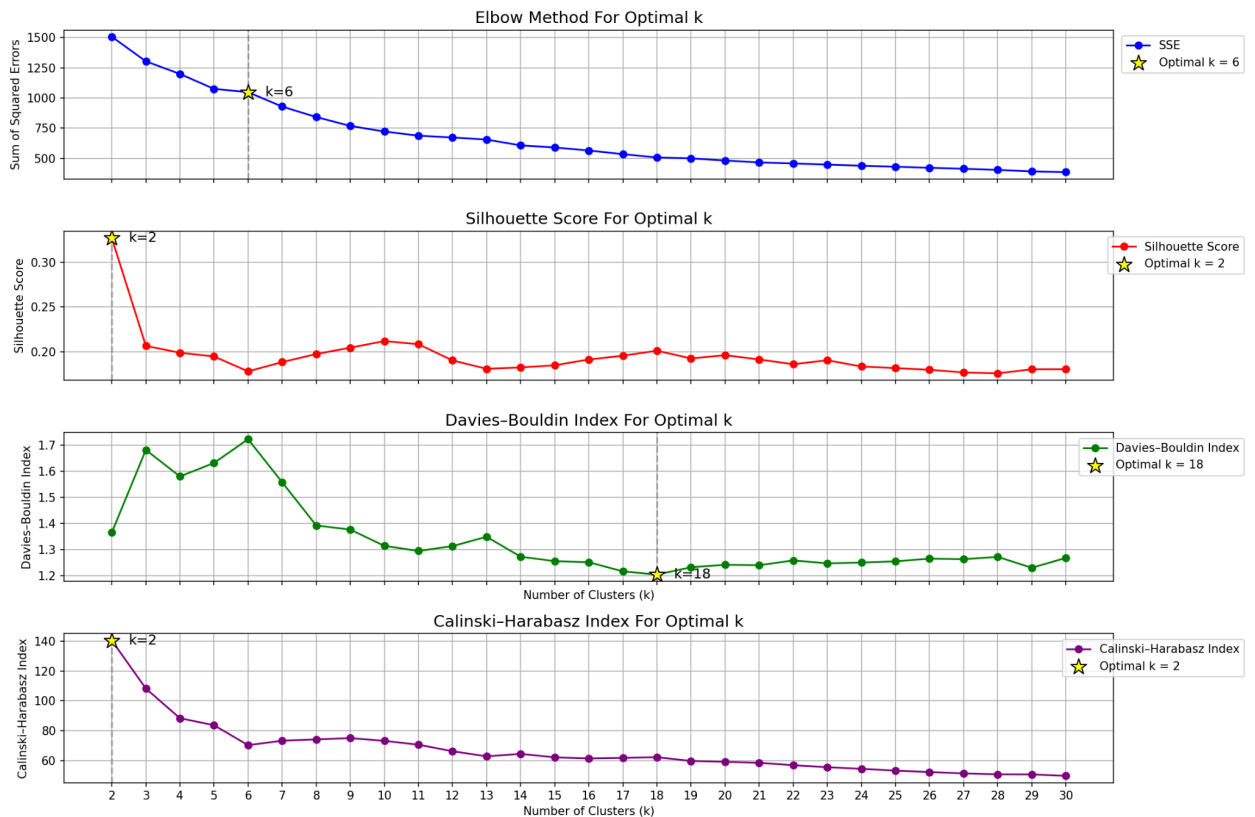


**Figure 1. Comparison of Random Forest Predictions with Different Feature Sets**

As illustrated in Figure 1, each subplot compares the model’s predicted daily expenses (y-axis) to the actual values (x-axis), with the diagonal dashed line representing perfect predictions. Feature Set A exhibits a relatively wide scatter around this line, reflecting higher prediction errors. Feature Set B shows a slightly narrower spread, while Feature Set C appears even more tightly clustered around the diagonal, indicating improved accuracy. Notably, Feature Set D demonstrates the closest alignment of points with the perfect prediction line, visually confirming that the addition of both the lagged feature and rolling average yields the best predictive performance among the four configurations.


## B. Unsupervised Learning Experiments

Unsupervised learning experiments using K-Means clustering and Gaussian Mixture Models revealed that the dataset naturally segmented into distinct groups; while the Elbow method suggested an optimal  $k$  of 6, alternative metrics provided varied insights, and GMM clustering further confirmed the presence of six distinct clusters with high confidence, each exhibiting unique spending profiles. These findings underscore that careful feature engineering—particularly the integration of temporal dynamics—significantly improves predictive accuracy and offers valuable insights into underlying spending behaviors.



**Figure 2. Optimal k Value Using Four Different Algorithms**

To determine the optimal value of  $k$ , I use four algorithms—Elbow Method, Silhouette Score, Davies-Bouldin Index, and Calinski-Harabasz—and get three unique values of  $k=2$ , 6, and 18. As depicted in Figure 3, the data has been projected onto two principal components for visualization, with each color indicating a cluster and black stars marking the cluster centroids. The left subplot illustrates the K-Means solution for  $k=6$  as suggested by the Elbow Method, yielding moderately sized, distinct clusters. In contrast, the middle subplot uses  $k=2$  from the Silhouette Score, resulting in just two large clusters that capture broad groupings of points. Meanwhile, the right subplot with  $k=18$  from the Davies-Bouldin Index partitions the data into numerous smaller clusters, reflecting a more granular segmentation. These differing cluster configurations underscore the inherent variability in selecting an optimal  $k$ , further emphasizing the importance of examining multiple clustering metrics to gain a comprehensive understanding of the dataset's structure.




**Figure 3. K-Means Clustering Solutions for  $k=2$ ,  $k=6$ , and  $k=18$**

Cluster	Count	Mean Daily Expense (NTD)	Std. Dev (NTD)
0	25	260.92	50.41
1	96	220.35	34.43
2	115	220.47	31.00
3	42	231.38	30.98
4	28	309.14	103.65
5	46	320.24	59.16

**Table 3. Key Statistics for Clusters (Elbow Method,  $k=6$ )**

Building on these clustering solutions and metrics, we now turn our attention to a more detailed analysis of the clusters—particularly those derived from the Elbow method with  $k=6$ . This in-depth examination allows us to better understand the underlying spending behaviors captured by each cluster. The clustering analysis reveals that the dataset naturally segments into distinct groups with unique spending profiles. For example, Cluster 2 is the largest with 115 samples and shows a consistent spending pattern, with a mean daily expense of approximately 220.47 NTD and a low standard deviation of 31.00 NTD. In contrast, Cluster 5, although smaller with 46 samples, has a significantly higher mean daily expense of around 320.24 NTD and a greater variability (std  $\approx$  59.16 NTD), indicating more heterogeneous spending behavior. Cluster 4 also stands out, with a high mean expense of 309.14 NTD, but its small size (28 samples) and very high standard deviation (103.65 NTD) suggest the presence of extreme spending values. These differences illustrate that spending behaviors differ markedly across segments, supporting the value of tailored financial strategies for each group. Table 3 highlights the distinct financial behaviors across clusters and reinforces the conclusion that incorporating unsupervised learning provides valuable insights into spending patterns.




**Figure 4. Fraction of Points Changing Clusters (From  $k$  to  $k+1$ )**

Beyond examining fixed clustering solutions, I also investigated how cluster assignments change as we increment the number of clusters from  $k$  to  $k+1$ . By computing the fraction of points that switch cluster membership between successive  $k$  values, we can gauge the stability of the clustering structure. In most cases, the fraction of points that change clusters is relatively low, indicating that the underlying groupings are robust. However, my analysis revealed that there are specific transitions—most notably from  $k=7$  to  $k=8$  and  $k=20$  to  $k=21$ —where nearly 80% of the points are reassigned. This abrupt change suggests that the data structure becomes highly unstable when attempting to over-segment the dataset, reinforcing the idea that the optimal number of clusters should

be chosen carefully to capture the inherent patterns without fragmenting them excessively.

Besides that, I also experimented by clustering the data using only the initial X months as the training set and then evaluated how well the learned cluster structure fit the remaining data. By varying the cutoff—using offsets of 2, 4, 6, 8, and 10 months—I examined whether the clusters derived from a smaller historical window remain representative when applied to future data. If the spending patterns are consistent over time, I expect the cluster distributions in both the training and subsequent test sets to be similar. However, noticeable differences in these distributions could indicate evolving consumer behavior or structural shifts in the data. This analysis not only serves to validate the robustness of the clustering approach but also provides insight into the temporal dynamics of daily expenses.



**Figure 5. Cluster Distributions for k=6 for Various Training Cutoffs**

When I plotted the cluster assignments for various offsets (2, 4, 6, 8, and 10 months) and different k values, several interesting patterns emerged. For smaller offsets, the training data is relatively limited and may not capture the full diversity of spending behaviors. As a result, the distribution of clusters in the test set can differ substantially from that in the



training set, indicating that additional months of data are necessary to learn stable, representative clusters. Conversely, for larger offsets, the training set grows and the resulting cluster distributions between training and test sets often become more aligned, suggesting that the learned clusters generalize better over time. Notably, when  $k=6$ , certain clusters remain fairly consistent across train and test splits, whereas others display significant shifts—particularly in smaller training windows—implying that new spending patterns or evolving behaviors appear in the latter months. These observations reinforce the importance of choosing an appropriate historical window when applying clustering models to time-dependent data, as an insufficient training period may overlook critical variations in consumer spending and lead to clusters that do not accurately represent future patterns.

Cluster	Count	Mean Daily Expense	Std. Dev (NTD)	Mean Confidence
0	49	268.43	54.05	0.99
1	99	220.91	34.25	0.98
2	132	223.13	31.51	0.97
3	22	222.86	31.15	0.98
4	3	442.00	276.22	0.95
5	47	328.32	52.77	0.99

**Table 4. GMM Cluster's Size and Key Statistics**

We also applied a Gaussian Mixture Model (GMM) to gain a probabilistic perspective on the data's cluster structure. Using six components (in line with the Elbow method's suggestion for K-Means), GMM produced clusters of varying sizes, with Cluster 2 being the largest (132 data points) and Cluster 4 containing only three data points. The presence of such a small cluster suggests that a handful of transactions were extreme outliers in terms of spending patterns. One advantage of GMM over K-Means is its soft assignment of points to clusters, reflected in each record's GMM\_Confidence value; most points displayed high confidence scores, indicating that they fit well within their assigned cluster's distribution. However, for points on the boundaries or in rare, high-expense categories (e.g., Cluster 4), the confidence may be lower or heavily skewed, highlighting the model's ability to capture nuanced variations in the data. Overall, the GMM approach confirms the existence of multiple spending profiles in the dataset, aligning with the K-Means results but offering additional insight into how firmly each transaction belongs to its respective cluster.

## **VI. Conclusion**

This experiments largely produced results that aligned with my expectations, while also uncovering some unexpected nuances. For instance, the improvement in predictive accuracy when incorporating both lag and rolling features confirmed that recent spending behavior plays a critical role in forecasting daily expenses. The clustering analyses further revealed distinct spending patterns—consistent with the heterogeneous nature of our simulated financial data—yet also highlighted sensitivity to outliers and shifts in cluster assignments when over-segmenting the dataset. Factors such as the variability inherent in daily expenditures, the presence of extreme values, and temporal influences (e.g., weekend effects, gym days) were found to significantly impact model performance. If more time were available, additional experiments could include testing advanced time-series models (like

LSTM networks) and exploring alternative clustering techniques such as hierarchical clustering to better capture non-linear and temporal dependencies. I also plan to use my real life spending notes if I have more time instead of generating the dataset. These experiments have underscored the importance of robust feature engineering and careful model selection, while raising further questions about the optimal balance between training window length and clustering stability, particularly in the face of rare, high-spending events.

## Appendix

*generate\_data.py*

```
import pandas as pd
import random
from datetime import datetime, timedelta, date

# Set random seed for reproducibility
random.seed(42)

# Helper function to generate a whole number with bias:
def biased_amount(min_val, max_val):
    """
    Generate a whole number between min_val and max_val (inclusive)
    with bias: numbers ending in 0 or 5 have higher probability.
    """
    values = list(range(min_val, max_val + 1))
    weights = [2 if (val % 10 == 0 or val % 10 == 5) else 1 for val in values]
    return random.choices(values, weights=weights, k=1)[0]

# Simulation period: one year (365 days)
start_date = date(2025, 1, 1)
num_days = 365
end_date = start_date + timedelta(days=num_days - 1)

records = []

# Create a list of month start dates for monthly records
months = pd.date_range(start=start_date, end=end_date, freq='MS')

# --- Revised Monthly Income ---

for m in months:
    m_date = m.date()

    # Scholarship: fixed 6000 NTD paid on the 20th of the same month.
    try:
        scholarship_date = date(m_date.year, m_date.month, 20)
```

```

except ValueError:
    # In case the month doesn't have 20 days (shouldn't occur)
    scholarship_date = date(m_date.year, m_date.month, 20)
    records.append({
        'Date': scholarship_date,
        'Category': 'Income',
        'Amount_NTD': 6000,
        'Description': f"Scholarship income for {m_date.strftime('%B %Y')}",
        'Payment_Method': None,
        'Time': None
    })

    # Parents: fixed 9000 NTD paid on a random day between 1st and 3rd of the
same month.
    parent_pay_day = random.randint(1, 3)
    parents_date = date(m_date.year, m_date.month, parent_pay_day)
    records.append({
        'Date': parents_date,
        'Category': 'Income',
        'Amount_NTD': 9000,
        'Description': f"Parents' income for {m_date.strftime('%B %Y')}",
        'Payment_Method': None,
        'Time': None
    })

    # Burger King income: paid on the 5th of the same month.
    bk_amount = biased_amount(1500, 4000)
    bk_pay_date = date(m_date.year, m_date.month, 5)
    records.append({
        'Date': bk_pay_date,
        'Category': 'Income',
        'Amount_NTD': bk_amount,
        'Description': f"Burger King income for {m_date.strftime('%B %Y')}",
        'Payment_Method': None,
        'Time': None
    })

    # Office Assistant income: paid on a random day between 25 and 30 of the next
month.

```

```

next_month_year = m_date.year
next_month = m_date.month + 1
if next_month > 12:
    next_month = 1
    next_month_year += 1
oa_pay_day = random.randint(25, 30)
try:
    oa_pay_date = date(next_month_year, next_month, oa_pay_day)
except ValueError:
    # Adjust if the day is invalid (e.g., February 30)
    oa_pay_date = date(next_month_year, next_month, 28)
# Only add office assistant income if it falls within the simulation period.
if start_date <= oa_pay_date <= end_date:
    oa_amount = biased_amount(4500, 5500)
    records.append({
        'Date': oa_pay_date,
        'Category': 'Income',
        'Amount_NTD': oa_amount,
        'Description': f"Office assistant income for {m_date.strftime('%B
%Y')}",
        'Payment_Method': None,
        'Time': None
    })

# Fixed monthly expenses (Transport, Entertainment, Groceries)
# Transport expense: between 200 and 500 NTD
records.append({
    'Date': m_date,
    'Category': 'Transport',
    'Amount_NTD': biased_amount(200, 500),
    'Description': 'Monthly transport expense',
    'Payment_Method': None,
    'Time': None
})

# Entertainment subscriptions: fixed monthly amounts
records.append({
    'Date': m_date,
    'Category': 'Entertainment',

```

```

        'Amount_NTD': 99,
        'Description': 'Spotify subscription (student)',
        'Payment_Method': None,
        'Time': None
    })
    records.append({
        'Date': m_date,
        'Category': 'Entertainment',
        'Amount_NTD': 600,
        'Description': 'ChatGPT Plus subscription',
        'Payment_Method': None,
        'Time': None
    })
    records.append({
        'Date': m_date,
        'Category': 'Entertainment',
        'Amount_NTD': 200,
        'Description': 'WuxiaWorld subscription',
        'Payment_Method': None,
        'Time': None
    })

# Groceries (household items): between 500 and 1000 NTD
records.append({
    'Date': m_date,
    'Category': 'Groceries',
    'Amount_NTD': biased_amount(500, 1000),
    'Description': 'Monthly groceries/household items',
    'Payment_Method': None,
    'Time': None
})

# Occasional special weekend events: one per semester (2 per year)
def get_random_weekend(start, end):
    """Return a random weekend date (Saturday or Sunday) between start and
    end."""
    delta = (end - start).days
    while True:
        rand_day = start + timedelta(days=random.randint(0, delta))

```

```

        if rand_day.weekday() >= 5:
            return rand_day

first_half = (start_date, date(start_date.year, 6, 30))
second_half = (date(start_date.year, 7, 1), end_date)
special_events = [
    {'Description': 'Korean BBQ with friends', 'Amount_NTD': 700},
    {'Description': 'Haidilao with friends', 'Amount_NTD': 500}
]
special_date1 = get_random_weekend(first_half[0], first_half[1])
special_date2 = get_random_weekend(second_half[0], second_half[1])
special_dates = [special_date1, special_date2]

for i, event in enumerate(special_events):
    records.append({
        'Date': special_dates[i],
        'Category': 'Food & Drink',
        'Amount_NTD': event['Amount_NTD'],
        'Description': event['Description'],
        'Payment_Method': None,
        'Time': 'Evening'
    })

# Gym supplements
protein_months = random.sample(list(months), 2)
for m in protein_months:
    records.append({
        'Date': m.date(),
        'Category': 'Gym',
        'Amount_NTD': 1500,
        'Description': 'Protein powder purchase (2.5kg)',
        'Payment_Method': None,
        'Time': None
    })

creatine_months = random.sample(list(months), 4)
for m in creatine_months:
    records.append({
        'Date': m.date(),
        'Category': 'Gym',

```



```

        'Amount_NTD': 600,
        'Description': 'Creatine purchase',
        'Payment_Method': None,
        'Time': None
    })

# Define gym days: roughly 3 days per week
gym_days = set()
for day in range(num_days):
    current_date = start_date + timedelta(days=day)
    if random.random() < 3/7:
        gym_days.add(current_date)

# Simulate daily expenses for meals, coffee, and laundry.
for day in range(num_days):
    current_date = start_date + timedelta(days=day)
    weekday = current_date.weekday() # Monday=0, Sunday=6

    # Lunch at around 12:00 PM (Meal expense: 80-105 NTD)
    lunch_time = "12:00"
    lunch_amount = biased_amount(80, 105)
    lunch_payment = 'Cash' if random.random() < 0.7 else 'Card'
    records.append({
        'Date': current_date,
        'Category': 'Meal',
        'Amount_NTD': lunch_amount,
        'Description': 'Lunch meal',
        'Payment_Method': lunch_payment,
        'Time': lunch_time
    })

# Coffee: approximately 3 times per week around 12:20-13:00 PM (40-60 NTD)
if random.random() < (3/7):
    coffee_time = f"12:{random.randint(20, 59):02d}"
    coffee_amount = biased_amount(40, 60)
    coffee_payment = 'Cash' if random.random() < 0.7 else 'Card'
    records.append({
        'Date': current_date,
        'Category': 'Coffee',

```

```

        'Amount_NTD': coffee_amount,
        'Description': 'Coffee purchase',
        'Payment_Method': coffee_payment,
        'Time': coffee_time
    })

# Dinner: between 5 PM and 8 PM.
# On gym days, dinner is more expensive (100-150 NTD);
# on weekends, moderately higher (100-250 NTD);
# otherwise, normal dinner: 80-110 NTD.
dinner_hour = random.randint(17, 20)
dinner_minute = random.randint(0, 59)
dinner_time = f"{dinner_hour}:{dinner_minute:02d}"
if current_date in gym_days:
    dinner_amount = biased_amount(100, 150)
elif weekday >= 5: # weekend
    dinner_amount = biased_amount(100, 250)
else:
    dinner_amount = biased_amount(80, 110)
dinner_payment = 'Cash' if random.random() < 0.7 else 'Card'
records.append({
    'Date': current_date,
    'Category': 'Meal',
    'Amount_NTD': dinner_amount,
    'Description': 'Dinner meal',
    'Payment_Method': dinner_payment,
    'Time': dinner_time
})

# Laundry: 60 NTD per week on Saturday
if current_date.weekday() == 5:
    records.append({
        'Date': current_date,
        'Category': 'Laundry',
        'Amount_NTD': 60,
        'Description': 'Weekly laundry expense',
        'Payment_Method': None,
        'Time': None
    })

```

```
# Create DataFrame and sort by Date and Time
df = pd.DataFrame(records)
df['Date'] = pd.to_datetime(df['Date'])
df.sort_values(by=['Date', 'Time'], inplace=True)

# Save the dataset to a CSV file and preview the first 20 records.
df.to_csv('daily_expenses.csv', index=False)
print("Data generation complete. Preview of the first 20 records:")
print(df.head(20))
```

*training.ipynb*

```
In [23]: import pandas as pd
import numpy as np

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import cross_val_score

import matplotlib.pyplot as plt
from matplotlib.lines import Line2D
```

```
In [24]: # Load the dataset with Date parsed as datetime
df = pd.read_csv('data\\daily_expenses.csv', parse_dates=['Date'])
```

```
In [25]: # Exclude 'Income' to focus on daily expenses
expense_df = df[df['Category'] != 'Income']

# Group by date to get total daily expense
daily_expense = expense_df.groupby('Date')['Amount_NTD'].sum().reset_index()
daily_expense.rename(columns={'Amount_NTD': 'DailyExpense'}, inplace=True)

# Sort by date (important if you want rolling features)
daily_expense.sort_values('Date', inplace=True)

# Drop daily expenses larger than 1000 NTD
daily_expense = daily_expense[daily_expense['DailyExpense'] <= 1000]
```

## Feature Engineering

```
In [26]: # (a) Day of week (Monday=0, Sunday=6)
daily_expense['DayOfWeek'] = daily_expense['Date'].dt.dayofweek

# (b) Weekend indicator
daily_expense['IsWeekend'] = daily_expense['DayOfWeek'].isin([5, 6]).astype(int)

# (c) Month and Day of month
daily_expense['Month'] = daily_expense['Date'].dt.month
daily_expense['Day'] = daily_expense['Date'].dt.day

# (d) Lag features (Will be used later)
# Create a lag feature (previous day's expense)
daily_expense['Lag1'] = daily_expense['DailyExpense'].shift(1)
# Fill missing values (e.g., first day) with the mean or zero
daily_expense['Lag1'].fillna(daily_expense['DailyExpense'].mean(), inplace=True)

# (e) Rolling 7-day average of expenses (Will be used later)
# This can help the model learn from recent spending trends
daily_expense['Rolling7'] = (
    daily_expense['DailyExpense']
    .rolling(window=7, min_periods=1)
    .mean()
)

# (f) Log transform the target to reduce the impact of large spikes
daily_expense['LogExpense'] = np.log1p(daily_expense['DailyExpense'])
```

# Random Forest

## Base Function

```
In [27]: # Function to train model and evaluate RMSE for a given feature set
def evaluate_feature_set(X, y, tscv=None, use_tscv=True):
    # Split data into training and test sets
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )

    # Set up a simple parameter grid (or use your existing grid)
    rf = RandomForestRegressor(random_state=42)
    param_grid = {
        'n_estimators': [100, 200, 300, 400, 500],
        'max_depth': [10, 20, 30, 40, 50],
        'min_samples_split': [2, 5, 10, 15, 20],
        'min_samples_leaf': [1, 2, 4, 8, 16],
        'max_features': ['sqrt', 'log2'],
        'bootstrap': [True, False]
    }

    if use_tscv:
        grid_search = GridSearchCV(
            rf, param_grid, cv=tscv, scoring='neg_mean_squared_error', n_jobs=-1
        )
    else:
        grid_search = GridSearchCV(
            rf, param_grid, cv=5, scoring='neg_mean_squared_error', n_jobs=-1
        )

    grid_search.fit(X_train, y_train)
```

```

print("Best Params from Grid Search:", grid_search.best_params_)

best_rf = grid_search.best_estimator_
y_pred_log = best_rf.predict(X_test)
y_pred = np.expm1(y_pred_log)
y_test_exp = np.expm1(y_test)

mse = mean_squared_error(y_test_exp, y_pred)
rmse = np.sqrt(mse)
return {'rmse': rmse, 'grid_search': grid_search, 'y_pred': y_pred, 'y_test': y_test_exp}

```

## Evaluate

### Without Lag1 and Rolling7

```

In [28]: tscv = TimeSeriesSplit(n_splits=5)

features_A = ['DayOfWeek', 'IsWeekend', 'Month', 'Day'] # without Lag1 and Rolling7

X = daily_expense[features_A]
y = daily_expense['LogExpense']

result = evaluate_feature_set(X, y, tscv, use_tscv=True)
rmse_A, best_params_A = result['rmse'], result['grid_search'].best_params_
best_rf_A = result['grid_search'].best_estimator_
y_pred_A, y_test_A = result['y_pred'], result['y_test']
feature_importance_A = result['grid_search'].best_estimator_.feature_importances_

print("Feature Set A (no Lag1, no Rolling7) RMSE:", rmse_A)

```

Best Params from Grid Search: {'bootstrap': True, 'max\_depth': 20, 'max\_features': 'sqrt', 'min\_samples\_leaf': 4, 'min\_samples\_split': 2, 'n\_estimators': 100}

Feature Set A (no Lag1, no Rolling7) RMSE: 38.66275658792042

### With Lag1

```

In [29]: tscv = TimeSeriesSplit(n_splits=5)

features_B = ['DayOfWeek', 'IsWeekend', 'Month', 'Day', 'Lag1'] # with Lag1 only

X = daily_expense[features_B]
y = daily_expense['LogExpense']

result = evaluate_feature_set(X, y, tscv, use_tscv=True)
rmse_B, best_params_B = result['rmse'], result['grid_search'].best_params_
best_rf_B = result['grid_search'].best_estimator_
y_pred_B, y_test_B = result['y_pred'], result['y_test']
feature_importance_B = result['grid_search'].best_estimator_.feature_importances_

print("Feature Set B (with Lag1, no Rolling7) RMSE:", rmse_B)

```

Best Params from Grid Search: {'bootstrap': True, 'max\_depth': 10, 'max\_features': 'sqrt', 'min\_samples\_leaf': 8, 'min\_samples\_split': 2, 'n\_estimators': 100}

Feature Set B (with Lag1, no Rolling7) RMSE: 39.3385956464684

### With Rolling7

```

In [30]: tscv = TimeSeriesSplit(n_splits=5)

features_C = ['DayOfWeek', 'IsWeekend', 'Month', 'Day', 'Rolling7'] # with Rolling7 only

X = daily_expense[features_C]
y = daily_expense['LogExpense']

result = evaluate_feature_set(X, y, tscv, use_tscv=True)
rmse_C, best_params_C = result['rmse'], result['grid_search'].best_params_
best_rf_C = result['grid_search'].best_estimator_
y_pred_C, y_test_C = result['y_pred'], result['y_test']
feature_importance_C = result['grid_search'].best_estimator_.feature_importances_

print("Feature Set C (no Lag1, with Rolling7) RMSE:", rmse_C)

```

Best Params from Grid Search: {'bootstrap': True, 'max\_depth': 10, 'max\_features': 'sqrt', 'min\_samples\_leaf': 4, 'min\_samples\_split': 2, 'n\_estimators': 100}

Feature Set C (no Lag1, with Rolling7) RMSE: 37.275187253048614

### With Lag1 & Rolling7

```

In [31]: tscv = TimeSeriesSplit(n_splits=5)

features_D = ['DayOfWeek', 'IsWeekend', 'Month', 'Day', 'Lag1', 'Rolling7'] # with both Lag1 and Rolling7

X = daily_expense[features_D]
y = daily_expense['LogExpense']

result = evaluate_feature_set(X, y, tscv, use_tscv=True)
rmse_D, best_params_D = result['rmse'], result['grid_search'].best_params_
best_rf_D = result['grid_search'].best_estimator_
y_pred_D, y_test_D = result['y_pred'], result['y_test']
feature_importance_D = result['grid_search'].best_estimator_.feature_importances_

print("Feature Set D (with Lag1, with Rolling7) RMSE:", rmse_D)

```

Best Params from Grid Search: {'bootstrap': True, 'max\_depth': 10, 'max\_features': 'sqrt', 'min\_samples\_leaf': 4, 'min\_samples\_split': 2, 'n\_estimators': 200}  
Feature Set D (with Lag1, with Rolling7) RMSE: 36.19773875656066

## Result

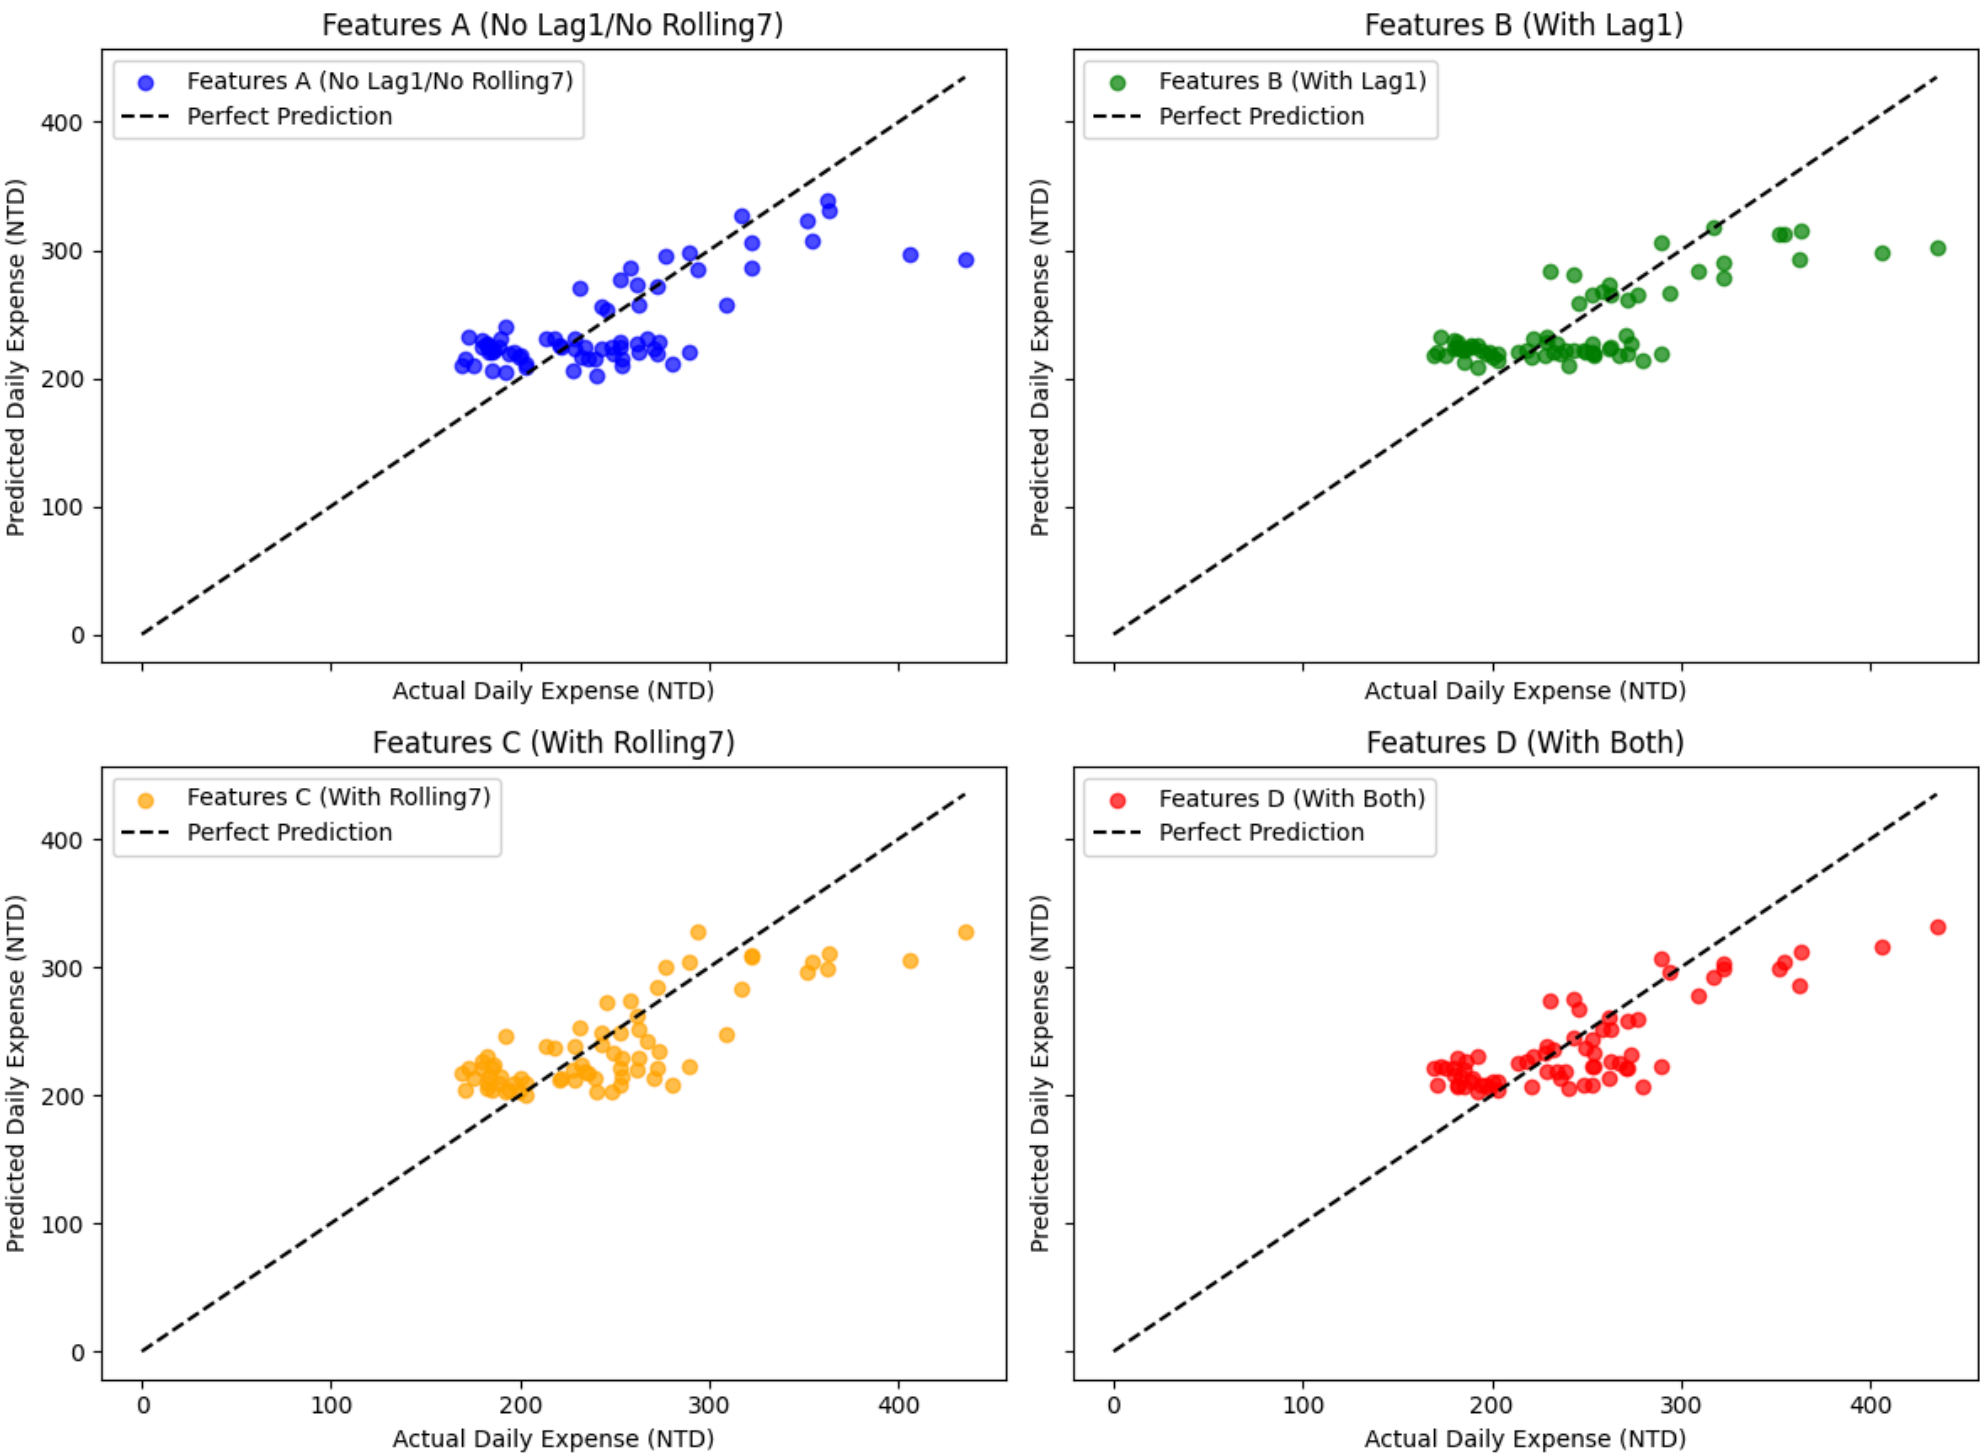
```
In [32]: # Create a 2x2 grid of subplots
fig, axes = plt.subplots(2, 2, figsize=(12, 10), sharex=True, sharey=True)
axes = axes.flatten()

# List of (true values, predictions, Label, color) for each feature set
plot_data = [
    (y_test_A, y_pred_A, 'Features A (No Lag1/No Rolling7)', 'blue'),
    (y_test_B, y_pred_B, 'Features B (With Lag1)', 'green'),
    (y_test_C, y_pred_C, 'Features C (With Rolling7)', 'orange'),
    (y_test_D, y_pred_D, 'Features D (With Both)', 'red'),
]

for ax, (y_test, y_pred, title, color) in zip(axes, plot_data):
    ax.scatter(y_test, y_pred, color=color, alpha=0.7, label=title)
    # Determine maximum value for perfect prediction line in each subplot
    max_val = max(y_test.max(), y_pred.max())
    ax.plot([0, max_val], [0, max_val], 'k--', label='Perfect Prediction')
    ax.set_xlabel('Actual Daily Expense (NTD)')
    ax.set_ylabel('Predicted Daily Expense (NTD)')
    ax.set_title(title)
    ax.legend()

fig.suptitle('Comparison of Random Forest Predictions with Different Feature Sets', fontsize=16)
plt.savefig('data\\experiment_result\\random_forest_comparison.png')
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()
```

Comparison of Random Forest Predictions with Different Feature Sets



## Feature Importance

```
In [33]: # Feature Set A: Without Lag1 and Rolling7
feature_importance_df_A = pd.DataFrame({'Feature': features_A, 'Importance': feature_importance_A})
print("Feature Importances for Feature Set A (No Lag1/No Rolling7):")
print(feature_importance_df_A)
print("\n" + "="*50 + "\n")

# Feature Set B: With Lag1 only
feature_importance_df_B = pd.DataFrame({'Feature': features_B, 'Importance': feature_importance_B})
print("Feature Importances for Feature Set B (With Lag1):")
```

```

print(feature_importance_df_B)
print("\n" + "="*50 + "\n")

# Feature Set C: With Rolling7 only
feature_importance_df_C = pd.DataFrame({'Feature': features_C, 'Importance': feature_importance_C})
print("Feature Importances for Feature Set C (With Rolling7):")
print(feature_importance_df_C)
print("\n" + "="*50 + "\n")

# Feature Set D: With both Lag1 and Rolling7
feature_importance_df_D = pd.DataFrame({'Feature': features_D, 'Importance': feature_importance_D})
print("Feature Importances for Feature Set D (With Both Lag1 & Rolling7):")
print(feature_importance_df_D)

```

Feature Importances for Feature Set A (No Lag1/No Rolling7):

	Feature	Importance
0	DayOfWeek	0.406320
1	IsWeekend	0.282965
2	Month	0.144396
3	Day	0.166320

=====

Feature Importances for Feature Set B (With Lag1):

	Feature	Importance
0	DayOfWeek	0.406350
1	IsWeekend	0.314623
2	Month	0.064919
3	Day	0.073103
4	Lag1	0.141005

=====

Feature Importances for Feature Set C (With Rolling7):

	Feature	Importance
0	DayOfWeek	0.335043
1	IsWeekend	0.239466
2	Month	0.078241
3	Day	0.108969
4	Rolling7	0.238281

=====

Feature Importances for Feature Set D (With Both Lag1 & Rolling7):

	Feature	Importance
0	DayOfWeek	0.292754
1	IsWeekend	0.229379
2	Month	0.054489
3	Day	0.079247
4	Lag1	0.130432
5	Rolling7	0.213698

## Cross Validation

In [34]: *# Function to perform cross-validation for a given feature set and best estimator*

```

def cross_val_feature_set(features, best_rf):
    X_set = daily_expense[features]
    y_set = daily_expense['LogExpense']

    cv_scores = cross_val_score(
        best_rf,
        X_set,
        y_set,
        cv=5,
        scoring='neg_mean_squared_error',
        n_jobs=-1
    )
    mse_scores = -cv_scores
    rmse_scores = np.sqrt(mse_scores)
    return rmse_scores, rmse_scores.mean(), rmse_scores.std()

```

In [35]: *# Perform cross-validation on each feature set using the best estimator from grid search*

```

rmse_scores_A, mean_rmse_A, std_rmse_A = cross_val_feature_set(features_A, best_rf_A)
rmse_scores_B, mean_rmse_B, std_rmse_B = cross_val_feature_set(features_B, best_rf_B)
rmse_scores_C, mean_rmse_C, std_rmse_C = cross_val_feature_set(features_C, best_rf_C)
rmse_scores_D, mean_rmse_D, std_rmse_D = cross_val_feature_set(features_D, best_rf_D)

```

In [36]:

```

print("Cross-Validation RMSE for Feature Set A (No Lag1/No Rolling7):")
print(rmse_scores_A)
print("Mean RMSE:", mean_rmse_A)
print("Standard Deviation of RMSE:", std_rmse_A)
print("\n" + "="*50 + "\n")
print("Cross-Validation RMSE for Feature Set B (With Lag1):")
print(rmse_scores_B)
print("Mean RMSE:", mean_rmse_B)
print("Standard Deviation of RMSE:", std_rmse_B)
print("\n" + "="*50 + "\n")
print("Cross-Validation RMSE for Feature Set C (With Rolling7):")
print(rmse_scores_C)
print("Mean RMSE:", mean_rmse_C)
print("Standard Deviation of RMSE:", std_rmse_C)
print("\n" + "="*50 + "\n")
print("Cross-Validation RMSE for Feature Set D (With Both Lag1 & Rolling7):")
print(rmse_scores_D)
print("Mean RMSE:", mean_rmse_D)
print("Standard Deviation of RMSE:", std_rmse_D)

```



Cross-Validation RMSE for Feature Set A (No Lag1/No Rolling7):  
[0.14123133 0.18620571 0.15099621 0.19368341 0.16927219]  
Mean RMSE: 0.16827776916692772  
Standard Deviation of RMSE: 0.019990040177141778

=====

Cross-Validation RMSE for Feature Set B (With Lag1):  
[0.13739528 0.17689063 0.15231579 0.19879546 0.16225005]  
Mean RMSE: 0.16552944218569143  
Standard Deviation of RMSE: 0.02103611466426487

=====

Cross-Validation RMSE for Feature Set C (With Rolling7):  
[0.12797353 0.17395989 0.14638658 0.17412378 0.15168051]  
Mean RMSE: 0.15482485750497763  
Standard Deviation of RMSE: 0.017553818642204157

=====

Cross-Validation RMSE for Feature Set D (With Both Lag1 & Rolling7):  
[0.12514251 0.16964064 0.14900217 0.17844627 0.1477058 ]  
Mean RMSE: 0.1539874791836075  
Standard Deviation of RMSE: 0.018652610616576897

```
In [37]: result_df = pd.DataFrame({
    'feature_set': ['A', 'B', 'C', 'D'],
    'RMSE': [rmse_A, rmse_B, rmse_C, rmse_D],
    'best_params': [best_params_A, best_params_B, best_params_C, best_params_D],
    'best_rf': [best_rf_A, best_rf_B, best_rf_C, best_rf_D],
    'features': [features_A, features_B, features_C, features_D],
    'feature_importance': [feature_importance_A, feature_importance_B, feature_importance_C, feature_importance_D],
    'cross_validation_RMSE': [rmse_scores_A, rmse_scores_B, rmse_scores_C, rmse_scores_D],
    'cross_validation_RMSE_mean': [mean_rmse_A, mean_rmse_B, mean_rmse_C, mean_rmse_D],
    'cross_validation_RMSE_std': [std_rmse_A, std_rmse_B, std_rmse_C, std_rmse_D],
})

result_df.to_csv('data\\experiment_result\\random_forest_result.csv', index=False)
```

## K-Means Clustering

```
In [38]: from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score, davies_bouldin_score, calinski_harabasz_score
```

```
In [39]: # Selecting features for clustering.
# We include some of the engineered features that capture temporal and trend information.
cluster_features = ['DayOfWeek', 'IsWeekend', 'Month', 'Day', 'Rolling7', 'Lag1']
```

```
In [40]: # Create a subset of data for clustering.
X_cluster = daily_expense[cluster_features].copy()

# Standardize the features (recommended for K-Means).
scaler = StandardScaler()
X_cluster_scaled = scaler.fit_transform(X_cluster)
```

```
In [41]: # Determine the optimal number of clusters using the Elbow Method.
sse = [] # Sum of Squared Errors for each k
sil_scores = []
db_scores = []
ch_scores = []

k_values = range(2, 31)
for k in k_values:
    # Note: In scikit-learn 1.4+ use n_init='auto'
    kmeans = KMeans(n_clusters=k, random_state=42, n_init='auto')
    labels = kmeans.fit_predict(X_cluster_scaled)

    sil = silhouette_score(X_cluster_scaled, labels)
    db = davies_bouldin_score(X_cluster_scaled, labels)
    ch = calinski_harabasz_score(X_cluster_scaled, labels)

    sse.append(kmeans.inertia_)
    sil_scores.append(sil)
    db_scores.append(db)
    ch_scores.append(ch)
```

```
In [42]: # Get the corresponding SSE value for optimal_k_elbow:
optimal_k_elbow = 6
sse_optimal = sse[optimal_k_elbow - k_values[0]] # adjust index if k_values does not start at 1

# Compute optimal k based on Silhouette Score and Davies-Bouldin Index
optimal_k_sil = k_values[sil_scores.index(max(sil_scores))]
sil_optimal = max(sil_scores)

optimal_k_db = k_values[db_scores.index(min(db_scores))]
db_optimal = min(db_scores)

optimal_k_ch = k_values[np.argmax(ch_scores)]
ch_optimal_value = max(ch_scores)
```



```

print(f"Optimal k based on Elbow Method: {optimal_k_elbow} (SSE={sse_optimal:.2f})")
print(f"Optimal k based on Silhouette Score: {optimal_k_sil} (Score={sil_optimal:.2f})")
print(f"Optimal k based on Davies-Bouldin Index: {optimal_k_db} (Score={db_optimal:.2f})")
print(f"Optimal k based on Calinski-Harabasz Index: {optimal_k_ch} (Score={ch_optimal_value:.2f})")

```

```

Optimal k based on Elbow Method: 6 (SSE=1047.89)
Optimal k based on Silhouette Score: 2 (Score=0.33)
Optimal k based on Davies-Bouldin Index: 18 (Score=1.20)
Optimal k based on Calinski-Harabasz Index: 2 (Score=140.54)

```

```

In [43]: fig, axes = plt.subplots(4, 1, figsize=(15, 10), sharex=True)

# --- 1) Elbow Method (SSE) ---
axes[0].plot(k_values, sse, marker='o', color='b', label="SSE")
axes[0].scatter(
    optimal_k_elbow, sse_optimal,
    s=200, marker='*',
    color='yellow', edgecolors='black', linewidths=1,
    zorder=10, label=f'Optimal k = {optimal_k_elbow}'
)
axes[0].axvline(optimal_k_elbow, color='gray', linestyle='--', alpha=0.7)
axes[0].text(
    optimal_k_elbow + 0.5, sse_optimal,
    f"k={optimal_k_elbow}",
    fontsize=12,
    verticalalignment='center'
)
axes[0].set_ylabel("Sum of Squared Errors")
axes[0].set_title("Elbow Method For Optimal k", fontsize=14)
axes[0].grid(True)
axes[0].legend(loc='upper right', bbox_to_anchor=(1.13, 1.0))

# --- 2) Silhouette Score ---
axes[1].plot(k_values, sil_scores, marker='o', color='r', label="Silhouette Score")
axes[1].scatter(
    optimal_k_sil, sil_optimal,
    s=200, marker='*',
    color='yellow', edgecolors='black', linewidths=1,
    zorder=10, label=f'Optimal k = {optimal_k_sil}'
)
axes[1].axvline(optimal_k_sil, color='gray', linestyle='--', alpha=0.7)
axes[1].text(
    optimal_k_sil + 0.5, sil_optimal,
    f"k={optimal_k_sil}",
    fontsize=12,
    verticalalignment='center'
)
axes[1].set_ylabel("Silhouette Score")
axes[1].set_title("Silhouette Score For Optimal k", fontsize=14)
axes[1].grid(True)
axes[1].legend(loc='upper right', bbox_to_anchor=(1.13, 1.0))


# --- 3) Davies-Bouldin Index ---
axes[2].plot(k_values, db_scores, marker='o', color='g', label="Davies-Bouldin Index")
axes[2].scatter(
    optimal_k_db, db_optimal,
    s=200, marker='*',
    color='yellow', edgecolors='black', linewidths=1,
    zorder=10, label=f'Optimal k = {optimal_k_db}'
)
axes[2].axvline(optimal_k_db, color='gray', linestyle='--', alpha=0.7)
axes[2].text(
    optimal_k_db + 0.5, db_optimal,
    f"k={optimal_k_db}",
    fontsize=12,
    verticalalignment='center'
)
axes[2].set_xlabel("Number of Clusters (k)")
axes[2].set_ylabel("Davies-Bouldin Index")
axes[2].set_title("Davies-Bouldin Index For Optimal k", fontsize=14)
axes[2].grid(True)
axes[2].legend(loc='upper right', bbox_to_anchor=(1.13, 1.0))

axes[3].plot(k_values, ch_scores, marker='o', color='purple', label="Calinski-Harabasz Index")
axes[3].scatter(
    optimal_k_ch, ch_optimal_value,
    s=200, marker='*',
    color='yellow', edgecolors='black', linewidths=1,
    zorder=10, label=f'Optimal k = {optimal_k_ch}'
)
axes[3].axvline(optimal_k_ch, color='gray', linestyle='--', alpha=0.7)
axes[3].text(
    optimal_k_ch + 0.5, ch_optimal_value,
    f"k={optimal_k_ch}",
    fontsize=12,
    verticalalignment='center'
)
axes[3].set_xlabel("Number of Clusters (k)")
axes[3].set_ylabel("Calinski-Harabasz Index")
axes[3].set_title("Calinski-Harabasz Index For Optimal k", fontsize=14)
axes[3].grid(True)
axes[3].legend(loc='upper right', bbox_to_anchor=(1.13, 1.0))

# Set common x-ticks across subplots
axes[3].set_xticks(k_values)

```

```
plt.tight_layout()
plt.savefig('data/experiment_result/kmeans_comparison_chart.png', dpi=150)
plt.show()
```



In [58]:

```
# Based on the elbow plot, choose an optimal number of clusters.
k_methods = [
    (optimal_k_elbow, "Elbow Method"),
    (optimal_k_sil, "Silhouette Score"),
    (optimal_k_db, "Davies-Bouldin Index")
]

# X_cluster_scaled is your scaled data for clustering
# Perform PCA once for consistent axes
pca = PCA(n_components=2, random_state=42)
X_pca_all = pca.fit_transform(X_cluster_scaled)

fig, axes = plt.subplots(1, 3, figsize=(25, 10), sharex=True, sharey=True)

for i, (k, method_name) in enumerate(k_methods):
    ax = axes[i]

    # Fit KMeans for the current k
    kmeans_model = KMeans(n_clusters=k, random_state=42, n_init='auto')
    clusters = kmeans_model.fit_predict(X_cluster_scaled)

    # Scatter plot of PCA-transformed data
    scatter = ax.scatter(
        X_pca_all[:, 0],
        X_pca_all[:, 1],
        c=clusters,
        cmap='viridis',
        alpha=0.7,
        s=50
    )

    # Plot cluster centroids in PCA space
    centers_2d = pca.transform(kmeans_model.cluster_centers_)
    ax.scatter(
        centers_2d[:, 0],
        centers_2d[:, 1],
        c='black',
        marker='*',
        s=200
    )

    # Title & Labels
    ax.set_title(f"K-Means (k={k}) from {method_name}", fontsize=13)
    ax.set_xlabel("Principal Component 1")
    if i == 0:
        ax.set_ylabel("Principal Component 2")

    # Add a colorbar for cluster labels
    cbar = plt.colorbar(scatter, ax=ax, fraction=0.046, pad=0.04)
    cbar.set_label('Cluster Label', rotation=270, labelpad=15)

    # Place a small text label in the top-left corner indicating the black star is for centroids
```

```


# Using Unicode star (U+2605) for a small star symbol
ax.text(
    0.02, 0.98, # x,y in Axes fraction
    u"\u2605 Centroids",
    transform=ax.transAxes,
    color='black',
    fontsize=12,
    ha='left',
    va='top',
    bbox=dict(facecolor='white', alpha=0.5, edgecolor='none') # optional background box
)

```

```

plt.tight_layout()
plt.savefig('data\\experiment_result\\kmeans_pca_comparison.png', dpi=200, bbox_inches='tight')
plt.show()

```



```

In [45]: features_to_describe = ['DailyExpense', 'Rolling7', 'Lag1']

for k, model in k_methods:
    kmeans_model = KMeans(n_clusters=k, random_state=42, n_init='auto')
    clusters = kmeans_model.fit_predict(X_cluster_scaled)
    col_name = f'Cluster_{model}_{k}'

    print(f'\n{col_name}')
    # Print counts of samples per cluster
    print(pd.Series(clusters).value_counts())

    # Optionally add the cluster labels to your DataFrame for further analysis
    daily_expense[col_name] = clusters

    # Use .describe() to calculate count, mean, std, min, 25%, 50%, 75%, and max
    stats = daily_expense.groupby(col_name)[features_to_describe].describe().round(2)
    print("\nDescriptive statistics:")
    print(stats)

    # Alternatively, if you want to compute specific metrics using .agg():
    custom_stats = daily_expense.groupby(col_name)[features_to_describe].agg({
        'DailyExpense': ['mean', 'std', 'min', 'max', lambda x: x.quantile(0.25), 'median', lambda x: x.quantile(0.75)],
        'Rolling7': ['mean', 'std', 'min', 'max', lambda x: x.quantile(0.25), 'median', lambda x: x.quantile(0.75)],
        'Lag1': ['mean', 'std', 'min', 'max', lambda x: x.quantile(0.25), 'median', lambda x: x.quantile(0.75)]
    })
    # Rename the lambda columns for clarity
    custom_stats.columns = ['_'.join(col).strip() for col in custom_stats.columns.values]
    print("\nCustom statistics:")
    print(custom_stats)

```

Cluster\_Elbow Method\_6  
2 115  
1 96  
5 46  
3 42  
4 28  
0 25  
Name: count, dtype: int64

Descriptive statistics:

Cluster_Elbow Method_6	DailyExpense						\
	count	mean	std	min	25%	50%	
0	25.0	260.92	50.41	195.0	218.00	251.0	
1	96.0	220.35	34.43	168.0	190.75	221.0	
2	115.0	220.47	31.00	168.0	193.50	221.0	
3	42.0	231.38	30.98	170.0	210.00	239.5	
4	28.0	309.14	103.65	195.0	256.75	296.5	
5	46.0	320.24	59.16	191.0	283.00	316.0	

Cluster_Elbow Method_6	Rolling7						\
	75%	max	count	mean	...	75%	max
0	298.00	396.0	25.0	245.36	...	253.29	270.14
1	242.25	304.0	96.0	239.99	...	249.43	266.29
2	240.50	303.0	115.0	237.07	...	245.79	262.57
3	249.75	289.0	42.0	270.13	...	271.07	323.71
4	322.75	759.0	28.0	254.42	...	261.43	320.57
5	354.00	451.0	46.0	242.07	...	252.04	284.14

Cluster_Elbow Method_6	Lag1							\
	count	mean	std	min	25%	50%	75%	
0	25.0	334.76	47.95	272.0	300.00	322.0	354.00	
1	96.0	224.66	37.49	168.0	194.00	219.0	240.50	
2	115.0	228.25	37.86	168.0	196.50	230.0	250.00	
3	42.0	266.19	45.20	185.0	242.25	253.5	282.25	
4	28.0	323.96	104.11	193.0	257.25	299.0	354.50	
5	46.0	210.54	28.72	169.0	187.25	200.5	239.00	

Cluster_Elbow Method_6	max	
0	451.0	
1	369.0	
2	377.0	
3	403.0	
4	759.0	
5	268.0	

[6 rows x 24 columns]

Custom statistics:

Cluster_Elbow Method_6	DailyExpense_mean			DailyExpense_std			DailyExpense_min			\
0	260.92	0000		50.41	4052				195	
1	220.35	4167		34.43	4065				168	
2	220.46	9565		31.00	2072				168	
3	231.38	0952		30.97	6744				170	
4	309.14	2857		103.65	3166				195	
5	320.23	9130		59.15	9740				191	

Cluster_Elbow Method_6	DailyExpense_max		DailyExpense_<lambda_0>		\
0	396		218.00		
1	304		190.75		
2	303		193.50		
3	289		210.00		
4	759		256.75		
5	451		283.00		

Cluster_Elbow Method_6	DailyExpense_median		DailyExpense_<lambda_1>		\
0	251.0		298.00		
1	221.0		242.25		
2	221.0		240.50		
3	239.5		249.75		
4	296.5		322.75		
5	316.0		354.00		

Cluster_Elbow Method_6	Rolling7_mean		Rolling7_std		Rolling7_min		...		\
0	245.35	5714	12.41	0378	217.85	7143	...		
1	239.98	8095	12.77	9745	215.00	0000	...		
2	237.07	2588	12.79	4384	200.50	0000	...		
3	270.12	9252	20.26	1182	241.71	4286	...		
4	254.41	8367	23.36	1982	220.71	4286	...		
5	242.06	5217	18.05	6241	202.42	8571	...		

Cluster_Elbow Method_6	Rolling7_<lambda_0>		Rolling7_median		\
0	239.42	8571	247.71	4286	
1	229.92	8571	240.14	2857	
2	228.28	5714	238.42	8571	
3	260.03	5714	265.00	0000	
4	238.07	1429	253.50	0000	
5	231.03	5714	237.50	0000	

	Rolling7_<lambda_1>	Lag1_mean	Lag1_std	Lag1_min	\
Cluster_Elbow Method_6					
0	253.285714	334.760000	47.946394	272.0	
1	249.428571	224.656250	37.492688	168.0	
2	245.785714	228.249605	37.860765	168.0	
3	271.071429	266.190476	45.203193	185.0	
4	261.428571	323.964286	104.108733	193.0	
5	252.035714	210.543478	28.721426	169.0	

	Lag1_max	Lag1_<lambda_0>	Lag1_median	\
Cluster_Elbow Method_6				
0	451.0	300.00	322.0	
1	369.0	194.00	219.0	
2	377.0	196.50	230.0	
3	403.0	242.25	253.5	
4	759.0	257.25	299.0	
5	268.0	187.25	200.5	

	Lag1_<lambda_1>
Cluster_Elbow Method_6	
0	354.00
1	240.50
2	250.00
3	282.25
4	354.50
5	239.00

[6 rows x 21 columns]

Cluster\_Silhouette Score\_2

1 253

0 99

Name: count, dtype: int64

Descriptive statistics:

	DailyExpense						\
	count	mean	std	min	25%	50%	
Cluster_Silhouette Score_2							
0	99.0	302.12	76.10	191.0	252.0	297.0	
1	253.0	222.24	32.47	168.0	192.0	222.0	

	Rolling7						\
	75%	max	count	mean	...	75%	
Cluster_Silhouette Score_2					...		
0	326.0	759.0	99.0	246.39	...	254.79	
1	244.0	304.0	253.0	243.67	...	253.71	

	Lag1							\
	max	count	mean	std	min	25%	50%	
Cluster_Silhouette Score_2								
0	320.57	99.0	273.99	86.45	169.0	202.0	253.0	
1	323.71	253.0	233.18	41.61	168.0	201.0	231.0	

	75%	max
Cluster_Silhouette Score_2		
0	319.5	759.0
1	253.0	403.0

[2 rows x 24 columns]

Custom statistics:

	DailyExpense_mean	DailyExpense_std	\
Cluster_Silhouette Score_2			
0	302.121212	76.097316	
1	222.237154	32.474885	

	DailyExpense_min	DailyExpense_max	\
Cluster_Silhouette Score_2			
0	191	759	
1	168	304	

	DailyExpense_<lambda_0>	DailyExpense_median	\
Cluster_Silhouette Score_2			
0	252.0	297.0	
1	192.0	222.0	

	DailyExpense_<lambda_1>	Rolling7_mean	\
Cluster_Silhouette Score_2			
0	326.0	246.389971	
1	244.0	243.666535	

	Rolling7_std	Rolling7_min	...	\
Cluster_Silhouette Score_2			...	
0	19.111819	202.428571	...	
1	18.550536	200.500000	...	

	Rolling7_<lambda_0>	Rolling7_median	\
Cluster_Silhouette Score_2			
0	234.357143	245.571429	
1	231.571429	242.714286	

	Rolling7_<lambda_1>	Lag1_mean	Lag1_std	\
Cluster_Silhouette Score_2				
0	254.785714	273.989899	86.446928	
1	253.714286	233.184603	41.607929	

	Lag1_min	Lag1_max	Lag1_<lambda_0>	Lag1_median	\
Cluster_Silhouette Score_2					
0	169.0	759.0	202.0	253.0	
1	168.0	403.0	201.0	231.0	

	Lag1_<lambda_1>	
Cluster_Silhouette Score_2		
0	319.5	
1	253.0	

[2 rows x 21 columns]

Cluster_Davies-Bouldin Index_18	
8	41
6	34
4	32
11	26
14	24
15	20
12	20
10	20
7	19
17	18
1	18
2	18
0	15
3	15
5	15
16	10
9	6
13	1

Name: count, dtype: int64

Descriptive statistics:

	DailyExpense					\
	count	mean	std	min	25%	
Cluster_Davies-Bouldin Index_18						
0	15.0	282.87	61.97	203.0	233.50	
1	18.0	234.83	34.50	180.0	214.25	
2	18.0	234.83	28.21	170.0	231.00	
3	15.0	294.93	40.14	234.0	263.00	
4	32.0	222.78	28.97	171.0	202.00	
5	15.0	226.53	34.19	178.0	201.50	
6	34.0	222.82	34.13	168.0	193.00	
7	19.0	276.26	54.27	196.0	233.00	
8	41.0	212.34	33.34	168.0	189.00	
9	6.0	333.50	208.53	243.0	244.50	
10	20.0	281.75	49.56	195.0	250.50	
11	26.0	231.31	32.14	169.0	206.50	
12	20.0	214.50	30.09	179.0	190.00	
13	1.0	253.00	NaN	253.0	253.00	
14	24.0	221.29	34.76	173.0	187.25	
15	20.0	205.30	26.35	175.0	187.25	
16	10.0	366.20	55.09	290.0	325.25	
17	18.0	315.83	69.13	191.0	278.00	

	Rolling7					...	\
	50%	75%	max	count	mean	...	
Cluster_Davies-Bouldin Index_18						...	
0	280.0	320.50	400.0	15.0	248.71	...	
1	236.5	253.50	297.0	18.0	249.15	...	
2	234.5	248.75	278.0	18.0	257.21	...	
3	297.0	319.50	362.0	15.0	235.21	...	
4	230.0	241.50	280.0	32.0	236.32	...	
5	224.0	244.00	289.0	15.0	262.09	...	
6	226.0	243.25	303.0	34.0	240.13	...	
7	277.0	318.50	403.0	19.0	254.11	...	
8	201.0	223.00	304.0	41.0	233.30	...	
9	248.0	256.75	759.0	6.0	321.07	...	
10	279.0	312.50	377.0	20.0	231.76	...	
11	236.0	253.50	276.0	26.0	258.28	...	
12	211.0	227.75	289.0	20.0	247.88	...	
13	253.0	253.00	253.0	1.0	311.29	...	
14	225.0	244.00	288.0	24.0	236.31	...	
15	199.0	215.50	268.0	20.0	218.00	...	
16	370.0	405.00	436.0	10.0	267.54	...	
17	316.5	365.25	451.0	18.0	242.40	...	

	Lag1						\
	75%	max	count	mean	std	min	
Cluster_Davies-Bouldin Index_18							
0	253.93	269.14	15.0	318.27	41.56	236.0	
1	255.75	261.86	18.0	255.28	31.78	209.0	
2	268.32	278.00	18.0	256.17	31.02	218.0	
3	240.79	261.14	15.0	242.53	47.92	174.0	
4	242.18	250.86	32.0	238.34	27.86	175.0	
5	266.00	272.86	15.0	254.93	39.15	185.0	
6	248.86	262.57	34.0	211.32	27.78	168.0	
7	261.43	283.57	19.0	368.11	45.77	301.0	
8	239.43	253.29	41.0	209.32	27.94	168.0	
9	322.54	323.71	6.0	251.17	5.64	244.0	
10	240.42	249.86	20.0	235.55	39.09	185.0	
11	264.04	267.86	26.0	234.27	27.93	187.0	
12	252.64	271.14	20.0	306.50	46.95	245.0	
13	311.29	311.29	1.0	759.00	NaN	759.0	
14	243.36	253.71	24.0	213.46	30.76	173.0	

15	222.25	233.14	20.0	198.79	22.20	171.0
16	275.82	284.14	10.0	205.10	26.90	170.0
17	250.46	274.29	18.0	219.06	37.03	169.0

	25%	50%	75%	max
Cluster_Davies-Bouldin Index_18				
0	295.00	318.0	342.00	400.0
1	235.25	249.0	279.75	322.0
2	232.50	247.5	267.00	324.0
3	196.00	245.0	282.50	317.0
4	229.75	240.0	253.00	309.0
5	231.00	248.0	286.00	318.0
6	190.00	210.0	229.25	271.0
7	328.00	362.0	400.00	451.0
8	189.00	203.0	223.00	314.0
9	247.00	251.5	254.50	259.0
10	200.00	233.5	256.75	314.0
11	213.25	238.0	258.50	276.0
12	279.25	295.0	324.25	403.0
13	759.00	759.0	759.00	759.0
14	190.25	211.5	228.25	288.0
15	181.50	196.0	206.75	248.0
16	191.00	193.5	225.00	250.0
17	186.50	219.0	247.25	296.0

[18 rows x 24 columns]

Custom statistics:

	DailyExpense_mean	DailyExpense_std \
Cluster_Davies-Bouldin Index_18		
0	282.866667	61.974496
1	234.833333	34.496803
2	234.833333	28.213993
3	294.933333	40.143433
4	222.781250	28.974666
5	226.533333	34.188274
6	222.823529	34.129407
7	276.263158	54.267693
8	212.341463	33.336624
9	333.500000	208.532731
10	281.750000	49.558976
11	231.307692	32.138163
12	214.500000	30.091964
13	253.000000	NaN
14	221.291667	34.758145
15	205.300000	26.352070
16	366.200000	55.087405
17	315.833333	69.130099

	DailyExpense_min	DailyExpense_max \
Cluster_Davies-Bouldin Index_18		
0	203	400
1	180	297
2	170	278
3	234	362
4	171	280
5	178	289
6	168	303
7	196	403
8	168	304
9	243	759
10	195	377
11	169	276
12	179	289
13	253	253
14	173	288
15	175	268
16	290	436
17	191	451

	DailyExpense_<lambda_0>	DailyExpense_median \
Cluster_Davies-Bouldin Index_18		
0	233.50	280.0
1	214.25	236.5
2	231.00	234.5
3	263.00	297.0
4	202.00	230.0
5	201.50	224.0
6	193.00	226.0
7	233.00	277.0
8	189.00	201.0
9	244.50	248.0
10	250.50	279.0
11	206.50	236.0
12	190.00	211.0
13	253.00	253.0
14	187.25	225.0
15	187.25	199.0
16	325.25	370.0
17	278.00	316.5

	DailyExpense_<lambda_1>	Rolling7_mean \
Cluster_Davies-Bouldin Index_18		
0	320.50	248.714286
1	253.50	249.150794
2	248.75	257.214286

3	319.50	235.209524
4	241.50	236.316667
5	244.00	262.085714
6	243.25	240.126050
7	318.50	254.112782
8	223.00	233.299652
9	256.75	321.071429
10	312.50	231.758929
11	253.50	258.280220
12	227.75	247.878571
13	253.00	311.285714
14	244.00	236.309524
15	215.50	217.996429
16	405.00	267.542857
17	365.25	242.396825

	Rolling7_std	Rolling7_min	...	\
Cluster_Davies-Bouldin Index_18				
0	10.785748	230.714286	...	
1	9.507370	224.428571	...	
2	11.907532	242.714286	...	
3	11.739007	213.428571	...	
4	7.830816	223.000000	...	
5	6.108542	251.142857	...	
6	10.610945	217.857143	...	
7	14.499791	222.428571	...	
8	9.881636	215.000000	...	
9	2.110711	318.285714	...	
10	12.252890	202.428571	...	
11	6.475383	242.857143	...	
12	11.412703	229.571429	...	
13	NaN	311.285714	...	
14	9.853647	216.428571	...	
15	7.373439	200.500000	...	
16	10.467269	252.285714	...	
17	12.760943	221.142857	...	

	Rolling7_<lambda_0>	Rolling7_median	\
Cluster_Davies-Bouldin Index_18			
0	240.642857	247.857143	
1	246.107143	248.785714	
2	247.642857	253.500000	
3	230.357143	235.571429	
4	229.250000	237.285714	
5	258.857143	262.714286	
6	232.678571	239.857143	
7	247.071429	254.000000	
8	224.857143	234.714286	
9	319.500000	321.285714	
10	223.464286	231.928571	
11	254.642857	258.642857	
12	240.750000	244.714286	
13	311.285714	311.285714	
14	228.857143	235.500000	
15	214.535714	219.000000	
16	259.964286	267.071429	
17	234.892857	240.214286	

	Rolling7_<lambda_1>	Lag1_mean	Lag1_std	\
Cluster_Davies-Bouldin Index_18				
0	253.928571	318.266667	41.561429	
1	255.750000	255.277778	31.783839	
2	268.321429	256.166667	31.020392	
3	240.785714	242.533333	47.917886	
4	242.178571	238.343750	27.862114	
5	266.000000	254.933333	39.147097	
6	248.857143	211.323529	27.781311	
7	261.428571	368.105263	45.771042	
8	239.428571	209.317073	27.938718	
9	322.535714	251.166667	5.636193	
10	240.419643	235.550000	39.088934	
11	264.035714	234.269231	27.930711	
12	252.642857	306.500000	46.948460	
13	311.285714	759.000000	NaN	
14	243.357143	213.458333	30.764151	
15	222.250000	198.785227	22.202398	
16	275.821429	205.100000	26.904977	
17	250.464286	219.055556	37.025384	

	Lag1_min	Lag1_max	Lag1_<lambda_0>	\
Cluster_Davies-Bouldin Index_18				
0	236.0	400.0	295.00	
1	209.0	322.0	235.25	
2	218.0	324.0	232.50	
3	174.0	317.0	196.00	
4	175.0	309.0	229.75	
5	185.0	318.0	231.00	
6	168.0	271.0	190.00	
7	301.0	451.0	328.00	
8	168.0	314.0	189.00	
9	244.0	259.0	247.00	
10	185.0	314.0	200.00	
11	187.0	276.0	213.25	
12	245.0	403.0	279.25	
13	759.0	759.0	759.00	
14	173.0	288.0	190.25	
15	171.0	248.0	181.50	



16	170.0	250.0	191.00
17	169.0	296.0	186.50
	Lag1_median	Lag1_<lambda_1>	
Cluster_Davies-Bouldin	Index_18		
0	318.0	342.00	
1	249.0	279.75	
2	247.5	267.00	
3	245.0	282.50	
4	240.0	253.00	
5	248.0	286.00	
6	210.0	229.25	
7	362.0	400.00	
8	203.0	223.00	
9	251.5	254.50	
10	233.5	256.75	
11	238.0	258.50	
12	295.0	324.25	
13	759.0	759.00	
14	211.5	228.25	
15	196.0	206.75	
16	193.5	225.00	
17	219.0	247.25	

[18 rows x 21 columns]

## Comparing How Cluster Assignments Change from k to k+1

```
In [46]: assignment_changes = []
previous_labels = None
previous_k = None

for k in k_values:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init='auto')
    labels = kmeans.fit_predict(X_cluster_scaled)

    if previous_labels is not None:
        # Count how many labels differ
        changes = np.sum(labels != previous_labels)
        fraction_changed = changes / len(labels)
        assignment_changes.append((previous_k, k, fraction_changed))

    previous_labels = labels
    previous_k = k

# Convert to a DataFrame for easier reading
changes_df = pd.DataFrame(assignment_changes, columns=["k_old", "k_new", "fraction_changed"])
print("Fraction of points changing clusters when going from k to k+1:")
print(changes_df)
```

Fraction of points changing clusters when going from k to k+1:

	k_old	k_new	fraction_changed
0	2	3	0.352273
1	3	4	0.116477
2	4	5	0.116477
3	5	6	0.221591
4	6	7	0.301136
5	7	8	0.735795
6	8	9	0.187500
7	9	10	0.090909
8	10	11	0.116477
9	11	12	0.076705
10	12	13	0.150568
11	13	14	0.025568
12	14	15	0.142045
13	15	16	0.110795
14	16	17	0.036932
15	17	18	0.073864
16	18	19	0.079545
17	19	20	0.056818
18	20	21	0.792614
19	21	22	0.068182
20	22	23	0.076705
21	23	24	0.085227
22	24	25	0.105114
23	25	26	0.059659
24	26	27	0.062500
25	27	28	0.088068
26	28	29	0.065341
27	29	30	0.076705

```
In [47]: # Identify the row with the maximum fraction_changed
max_idx = changes_df['fraction_changed'].idxmax()
max_k = changes_df.loc[max_idx, 'k_new']
max_val = changes_df.loc[max_idx, 'fraction_changed']
print(f"Maximum change in cluster assignments occurs when going from k={max_k-1} to k={max_k} ({max_val:.2%} change)")
```

Maximum change in cluster assignments occurs when going from k=20 to k=21 (79.26% change)

```
In [48]: plt.figure(figsize=(10, 5))

# Create a bar chart
bars = plt.bar(
    changes_df['k_new'],
    changes_df['fraction_changed'],
    color='skyblue',
```

```


    edgecolor='black'
)

# Highlight the bar with the maximum fraction changed
for bar in bars:
    bar_center = bar.get_x() + bar.get_width()/2
    if np.isclose(bar_center, max_k, atol=0.1):
        bar.set_color('orange')
        bar.set_edgecolor('red')

# Annotate the highest bar
plt.text(
    max_k,
    max_val + 0.02, # Slightly above the bar
    f"{max_val:.2f}",
    ha='center',
    color='red',
    fontsize=12,
    fontweight='bold'
)

plt.title("Fraction of Points Changing Clusters\n(Going from k to k+1)", fontsize=14)
plt.xlabel("New Number of Clusters (k+1)")
plt.ylabel("Fraction of Points Changed")
plt.ylim(0, 1) # Fractions range from 0 to 1
plt.xticks(changes_df['k_new'])
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.savefig('data/experiment_result/cluster_assignment_changes.png', dpi=150)
plt.show()

```



## Cluster on First X Month, Then See How The Data Fits

```

In [49]: cutoff_date = pd.to_datetime("2025-06-01")
train_data = daily_expense[daily_expense['Date'] < cutoff_date].copy()
test_data = daily_expense[daily_expense['Date'] >= cutoff_date].copy()

# Scale them separately or together, depending on your approach
scaler = StandardScaler()
train_features = train_data[cluster_features]
test_features = test_data[cluster_features]

X_train_scaled = scaler.fit_transform(train_features)
X_test_scaled = scaler.transform(test_features)

k = 6 # or your chosen k
kmeans_split = KMeans(n_clusters=k, random_state=42, n_init='auto')
kmeans_split.fit(X_train_scaled)

# Assign clusters to training data
train_clusters = kmeans_split.predict(X_train_scaled)
train_data['Cluster'] = train_clusters

# Assign clusters to test data
test_clusters = kmeans_split.predict(X_test_scaled)
test_data['Cluster'] = test_clusters

print("Training cluster distribution:")
print(train_data['Cluster'].value_counts())

print("\nTest cluster distribution:")
print(test_data['Cluster'].value_counts())

```

```
Training cluster distribution:
Cluster
0      31
5      26
3      26
1      24
4      22
2      16
Name: count, dtype: int64
```

```
Test cluster distribution:
Cluster
4      115
3       47
1       23
2       22
Name: count, dtype: int64
```

## Check for 2, 4, 6, 8, 10 Months

```
In [50]: offsets = [2, 4, 6, 8, 10]
```

```
start_date = daily_expense['Date'].min()
split_results = {}
```

```
In [51]: # Loop over each offset to compute cluster assignments and count samples per cluster
for offset in offsets:
    cutoff_date = start_date + pd.DateOffset(months=offset)

    # Split the data
    train_data = daily_expense[daily_expense['Date'] < cutoff_date].copy()
    test_data = daily_expense[daily_expense['Date'] >= cutoff_date].copy()

    # Scale the features (fit scaler on train and transform test)
    scaler = StandardScaler()
    train_features = train_data[cluster_features]
    test_features = test_data[cluster_features]

    X_train_scaled = scaler.fit_transform(train_features)
    X_test_scaled = scaler.transform(test_features)

    # Fit K-Means on training data with chosen k (e.g., 6)
    k = 6
    kmeans_split = KMeans(n_clusters=k, random_state=42, n_init='auto')
    kmeans_split.fit(X_train_scaled)

    # Predict clusters for training and test data
    train_clusters = kmeans_split.predict(X_train_scaled)
    test_clusters = kmeans_split.predict(X_test_scaled)

    train_data['Cluster'] = train_clusters
    test_data['Cluster'] = test_clusters

    # Count samples per cluster
    train_counts = train_data['Cluster'].value_counts().sort_index()
    test_counts = test_data['Cluster'].value_counts().sort_index()

    # Store results for this offset
    split_results[offset] = {
        'cutoff_date': cutoff_date,
        'train_counts': train_counts,
        'test_counts': test_counts
    }
```

```
In [52]: fig, axes = plt.subplots(nrows=len(offsets), ncols=1, figsize=(8, 3 * len(offsets)), sharex=True)
```

```
x_positions = np.arange(k) # k is the number of clusters, e.g. 6
bar_width = 0.4
```

```
for i, offset in enumerate(offsets):
    result = split_results[offset]
    cutoff_date = result['cutoff_date']
    train_counts = result['train_counts']
    test_counts = result['test_counts']

    # Ensure both Series cover all clusters [0..k-1], fill missing with 0
    train_counts = train_counts.reindex(range(k), fill_value=0)
    test_counts = test_counts.reindex(range(k), fill_value=0)

    ax = axes[i] if len(offsets) > 1 else axes # handle single-subplot case


    # Plot grouped bars
    ax.bar(x_positions - bar_width/2, train_counts.values, width=bar_width,
           color='skyblue', edgecolor='black', label='Train')
    ax.bar(x_positions + bar_width/2, test_counts.values, width=bar_width,
           color='lightgreen', edgecolor='black', label='Test')

    ax.set_xticks(x_positions)
    ax.set_xticklabels([str(c) for c in range(k)])
    ax.set_ylabel("Count")
    ax.set_title(f"Offset {offset} months (Cutoff: {cutoff_date.date()})")
    ax.legend(loc='upper right')
```

```
plt.suptitle("Cluster Distributions (Grouped Bar) K=6 for Various Training Cutoffs", fontsize=14, y=1)
plt.savefig('data/experiment_result/train_test_cluster_distribution_k_6.png', dpi=150, bbox_inches='tight')
```

```
plt.tight_layout()
plt.show()
```

## Cluster Distributions (Grouped Bar) K=6 for Various Training Cutoffs



## Check for 2, 4, 6, 8, 10 Months Using K=6 and K=2

```
In [54]: offsets = [2, 4, 6, 8, 10]
k_vals = [2, 6]
start_date = daily_expense['Date'].min()

split_results = {offset: {} for offset in offsets}

In [55]: for offset in offsets:
    cutoff_date = start_date + pd.DateOffset(months=offset)

    # Split the data
    train_data = daily_expense[daily_expense['Date'] < cutoff_date].copy()
    test_data = daily_expense[daily_expense['Date'] >= cutoff_date].copy()

    # Scale the features (fit scaler on train, then transform test)
    scaler = StandardScaler()
    train_features = train_data[cluster_features]
    test_features = test_data[cluster_features]

    X_train_scaled = scaler.fit_transform(train_features)
    X_test_scaled = scaler.transform(test_features)

    for k in k_vals:
        # Fit K-Means on training data
        kmeans_split = KMeans(n_clusters=k, random_state=42, n_init='auto')
        kmeans_split.fit(X_train_scaled)

        # Predict clusters for training and test data
        train_clusters = kmeans_split.predict(X_train_scaled)
        test_clusters = kmeans_split.predict(X_test_scaled)

        # Count samples per cluster
        train_counts = pd.Series(train_clusters).value_counts().sort_index()
        test_counts = pd.Series(test_clusters).value_counts().sort_index()

        # Store the results
        split_results[offset][k] = {
            'cutoff_date': cutoff_date,
            'train_counts': train_counts,
            'test_counts': test_counts
        }
```

```
In [56]: n_offsets = len(offsets)
n_kvals = len(k_vals)
fig, axes = plt.subplots(
    nrows=n_offsets, ncols=n_kvals,
    figsize=(12, 4 * n_offsets),
    sharex=False, sharey=False
)

bar_width = 0.4

for i, offset in enumerate(offsets):
    for j, k in enumerate(k_vals):
        ax = axes[i, j] if n_offsets > 1 else axes[j] # handle single-row case
        results = split_results[offset][k]

        cutoff_date = results['cutoff_date']
        train_counts = results['train_counts']
        test_counts = results['test_counts']

        # Ensure we have a count for each cluster index [0..k-1]
        train_counts = train_counts.reindex(range(k), fill_value=0)
        test_counts = test_counts.reindex(range(k), fill_value=0)

        x_positions = np.arange(k)

        # Grouped bar chart: train vs test
        ax.bar(
            x_positions - bar_width/2,
            train_counts.values,
            width=bar_width,
            color='skyblue',
            edgecolor='black',
            label='Train'
        )
        ax.bar(
            x_positions + bar_width/2,
            test_counts.values,
            width=bar_width,
            color='lightgreen',
            edgecolor='black',
            label='Test'
        )

        ax.set_xticks(x_positions)
        ax.set_xticklabels([str(c) for c in range(k)])
        ax.set_ylabel("Count")
        ax.set_title(f"Offset {offset} mo (k={k}), Cutoff: {cutoff_date.date()}")


        if i == 0 and j == 0:
            ax.legend(loc='upper right')

plt.suptitle("Train/Test Cluster Distributions for Various Offsets & k Values", fontsize=14, y=1)
```


```
plt.savefig('data/experiment_result/train_test_cluster_distribution_multiple_offsets_and_k_vals.png', dpi=150, bbox_inches='tight')
plt.tight_layout()
plt.show()
```

# Train/Test Cluster Distributions for Various Offsets & k Values


Offset 2 mo (k=2), Cutoff: 2025-03-02




Offset 2 mo (k=6), Cutoff: 2025-03-02




Offset 4 mo (k=2), Cutoff: 2025-05-02




Offset 4 mo (k=6), Cutoff: 2025-05-02




Offset 6 mo (k=2), Cutoff: 2025-07-02




Offset 6 mo (k=6), Cutoff: 2025-07-02




Offset 8 mo (k=2), Cutoff: 2025-09-02




Offset 8 mo (k=6), Cutoff: 2025-09-02





Offset 10 mo (k=2), Cutoff: 2025-11-02



Offset 10 mo (k=6), Cutoff: 2025-11-02





# GMM (Gaussian Mixture Model)

```
In [ ]: from sklearn.mixture import GaussianMixture
from sklearn.preprocessing import StandardScaler
import seaborn as sns
```

```
In [ ]: cluster_features = ['DayOfWeek', 'IsWeekend', 'Month', 'Day', 'Rolling7', 'Lag1']
```

```
In [ ]: scaler = StandardScaler()
X_cluster_scaled = scaler.fit_transform(daily_expense[cluster_features])
```

```
In [ ]: n_components = 6
```

```
In [ ]: gmm = GaussianMixture(n_components=n_components, random_state=42)
gmm.fit(X_cluster_scaled)
```

Out [ ]:

▼ GaussianMixture

GaussianMixture(n\_components=6, random\_state=42)

```
In [ ]: # Predict cluster labels and get the membership probabilities
gmm_labels = gmm.predict(X_cluster_scaled)
gmm_probabilities = gmm.predict_proba(X_cluster_scaled)

# Add the GMM cluster assignments and probabilities to your DataFrame
daily_expense['GMM_Cluster'] = gmm_labels
# For example, you might store the maximum probability (confidence) for each point
daily_expense['GMM_Confidence'] = gmm_probabilities.max(axis=1)

print("GMM Cluster distribution:")
print(daily_expense['GMM_Cluster'].value_counts())

# Optionally, display the first few rows with cluster labels and confidence
print(daily_expense[['Date', 'DailyExpense', 'GMM_Cluster', 'GMM_Confidence']].head())
```

GMM Cluster distribution:

GMM\_Cluster

2	132
1	99
0	49
5	47
3	22
4	3

Name: count, dtype: int64

	Date	DailyExpense	GMM_Cluster	GMM_Confidence
1	2025-01-02	206	2	0.999993
2	2025-01-03	195	2	0.999993
3	2025-01-04	295	5	1.000000
4	2025-01-05	263	0	1.000000
5	2025-01-06	230	2	0.999994

## Descriptive Statistics

```
In [ ]: features_numeric = ['DailyExpense', 'Rolling7', 'Lag1']
print("=== Numeric Descriptive Statistics by Cluster ===")
profile_stats = daily_expense.groupby('GMM_Cluster')[features_numeric].describe().round(2)
print(profile_stats)
```



=== Numeric Descriptive Statistics by Cluster ===								
DailyExpense								
	count	mean	std	min	25%	50%	75%	max
GMM_Cluster								
0	49.0	268.43	54.05	191.0	222.00	258.0	309.00	403.0
1	99.0	220.91	34.25	169.0	191.00	221.0	243.00	304.0
2	132.0	223.13	31.51	168.0	194.75	227.0	246.25	303.0
3	22.0	222.86	31.15	168.0	200.50	217.5	245.75	289.0
4	3.0	442.00	276.22	253.0	283.50	314.0	536.50	759.0
5	47.0	328.32	52.77	245.0	292.50	321.0	358.00	451.0

Rolling7								
Lag1								
	count	mean	...	75%	max	count	mean	std
GMM_Cluster			...					
0	49.0	244.03	...	254.00	283.57	49.0	322.84	57.82
1	99.0	242.64	...	254.57	268.57	99.0	220.87	29.11
2	132.0	240.41	...	249.50	278.00	132.0	230.63	36.74
3	22.0	267.84	...	272.43	323.71	22.0	303.95	48.94
4	3.0	297.67	...	315.93	320.57	3.0	419.33	294.21
5	47.0	245.58	...	254.64	284.14	47.0	213.79	32.34

	min	25%	50%	75%	max
GMM_Cluster					
0	185.0	289.00	318.0	354.00	451.0
1	168.0	194.50	220.0	240.00	288.0
2	168.0	201.25	231.5	253.00	324.0
3	244.0	254.50	300.5	321.25	403.0
4	244.0	249.50	255.0	507.00	759.0
5	169.0	191.00	201.0	243.50	288.0

[6 rows x 24 columns]

```
In [ ]: # For categorical features, calculate percentages (e.g., DayOfWeek, IsWeekend)
categorical_features = ['DayOfWeek', 'IsWeekend', 'Month', 'Day']
print("\n=== Categorical Distribution by Cluster (in %): ===")
cat_profile = daily_expense.groupby('GMM_Cluster')[categorical_features].agg(
    lambda x: x.value_counts(normalize=True).mul(100).round(2).to_dict()
)
print(cat_profile)
```

=== Categorical Distribution by Cluster (in %): ===		
DayOfWeek		
GMM_Cluster		IsWeekend
0	{6: 100.0}	{1: 100.0}
1	{4: 24.24, 3: 23.23, 1: 21.21, 2: 20.2, 0: 11.11}	{0: 100.0}
2	{3: 20.45, 2: 20.45, 4: 19.7, 0: 19.7, 1: 19.7}	{0: 100.0}
3	{0: 59.09, 2: 18.18, 1: 13.64, 3: 4.55, 4: 4.55}	{0: 100.0}
4	{6: 66.67, 5: 33.33}	{1: 100.0}
5	{5: 100.0}	{1: 100.0}

Month	
GMM_Cluster	
0	{3: 10.2, 8: 10.2, 1: 8.16, 2: 8.16, 4: 8.16, ...
1	{11: 19.19, 12: 19.19, 9: 18.18, 8: 17.17, 10: ...
2	{1: 16.67, 3: 15.91, 4: 15.91, 2: 15.15, 5: 15...
3	{10: 36.36, 8: 13.64, 9: 13.64, 12: 13.64, 7: ...
4	{10: 66.67, 11: 33.33}
5	{5: 10.64, 8: 10.64, 1: 8.51, 4: 8.51, 6: 8.51...

Day	
GMM_Cluster	
0	{9: 6.12, 16: 6.12, 23: 6.12, 5: 4.08, 13: 4.0...
1	{4: 5.05, 10: 5.05, 24: 5.05, 25: 5.05, 30: 4...
2	{3: 4.55, 21: 4.55, 14: 4.55, 23: 3.79, 17: 3...
3	{22: 9.09, 2: 9.09, 13: 9.09, 3: 9.09, 16: 9.0...
4	{11: 33.33, 12: 33.33, 2: 33.33}
5	{8: 6.38, 15: 6.38, 4: 4.26, 12: 4.26, 20: 4.2...

```
In [ ]: # Investigate a small cluster (e.g., Cluster 4)
print("\n=== Detailed Profile for Cluster 4 ===")
cluster_4 = daily_expense[daily_expense['GMM_Cluster'] == 4]
print(cluster_4.describe())
```

	Date	DailyExpense	DayOfWeek	IsWeekend	Month
count	3	3.000000	3.000000	3.0	3.000000
mean	2025-10-18 16:00:00	442.000000	5.666667	1.0	10.333333
min	2025-10-11 00:00:00	253.000000	5.000000	1.0	10.000000
25%	2025-10-11 12:00:00	283.500000	5.500000	1.0	10.000000
50%	2025-10-12 00:00:00	314.000000	6.000000	1.0	10.000000
75%	2025-10-22 12:00:00	536.500000	6.000000	1.0	10.500000
max	2025-11-02 00:00:00	759.000000	6.000000	1.0	11.000000
std	NaN	276.219116	0.577350	0.0	0.577350

	Day	Lag1	Rolling7	LogExpense	GMM_Cluster \
count	3.000000	3.000000	3.000000	3.000000	3.0
mean	8.333333	419.333333	297.666667	5.974408	4.0
min	2.000000	244.000000	261.142857	5.537334	4.0
25%	6.500000	249.500000	286.214286	5.644953	4.0
50%	11.000000	255.000000	311.285714	5.752573	4.0
75%	11.500000	507.000000	315.928571	6.192946	4.0
max	12.000000	759.000000	320.571429	6.633318	4.0
std	5.507571	294.211375	31.969479	0.580692	0.0

	GMM_Confidence
count	3.000000e+00
mean	9.999999e-01
min	9.999997e-01
25%	9.999998e-01
50%	1.000000e+00
75%	1.000000e+00
max	1.000000e+00
std	1.893129e-07

The figure consists of three side-by-side boxplots, each representing a different variable: DailyExpense, Rolling7, and Lag1. Each plot has six clusters on the x-axis, labeled 0 through 5. The y-axis for each plot represents the value of the variable. The boxplots are color-coded: Cluster 0 is teal, Cluster 1 is yellow, Cluster 2 is purple, Cluster 3 is red, Cluster 4 is blue, and Cluster 5 is orange. The plots show the median, quartiles, and range (whiskers) of the data for each cluster. Outliers are present for Rolling7 (Cluster 3) and Lag1 (Cluster 0).

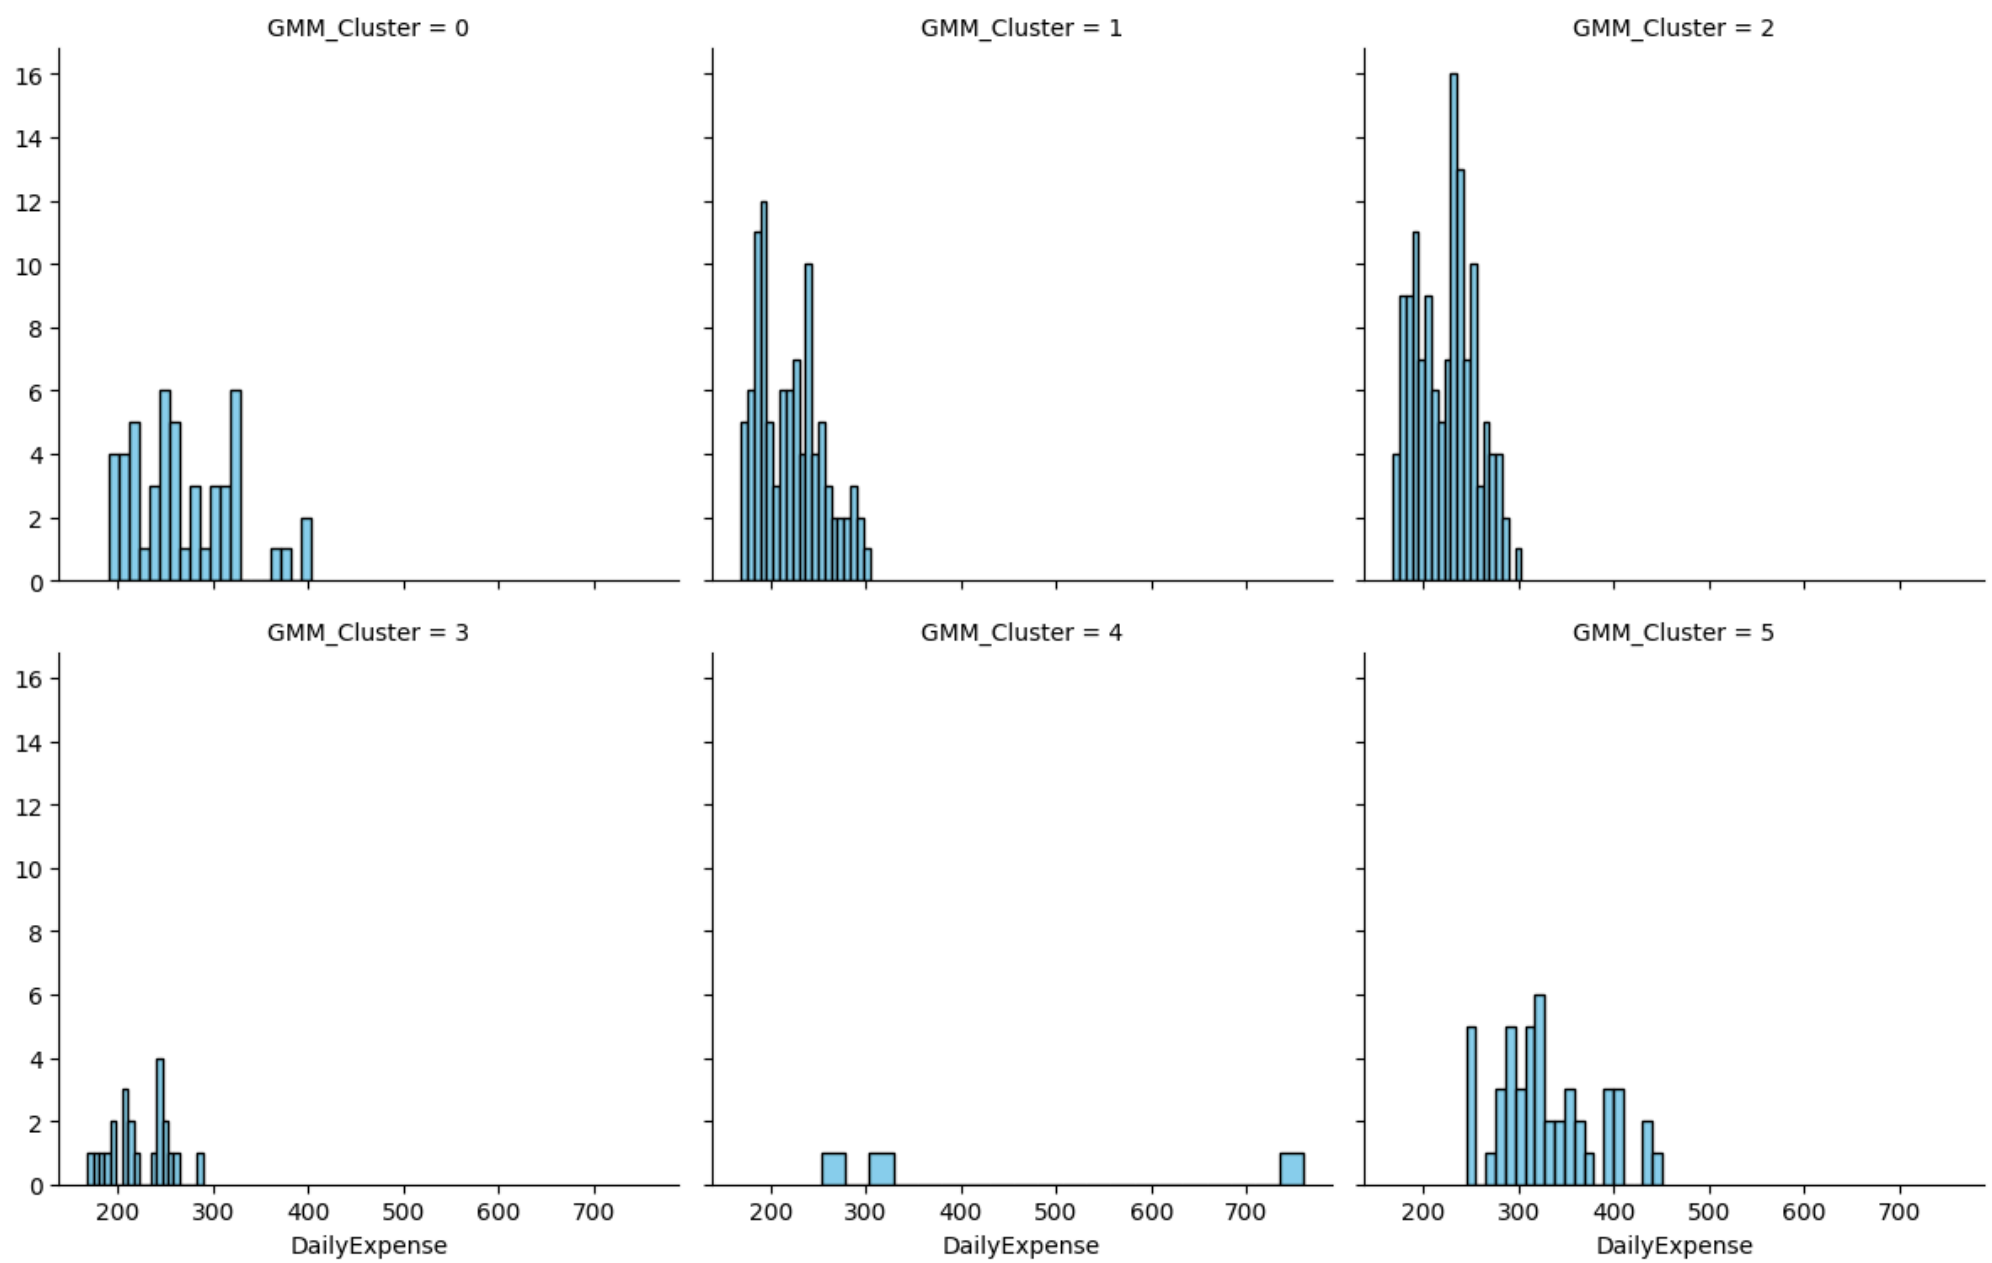
Cluster	DailyExpense (Median)	Rolling7 (Median)	Lag1 (Median)
0	~260	~245	~320
1	~220	~245	~220
2	~220	~240	~230
3	~220	~255	~300
4	~310	~310	~250
5	~330	~245	~200

The figure consists of three violin plots arranged horizontally, each showing the distribution of a different variable across six clusters (0 to 5). The y-axis for all plots ranges from -200 to 1200. The x-axis for all plots is labeled 'Cluster'.

- Violin Plot of DailyExpense by Cluster:** The y-axis is labeled 'DailyExpense'. Cluster 4 shows a very wide distribution with a median around 300 and a range from approximately -200 to 1200. Cluster 5 has a median around 300 and a range from approximately 200 to 500. Clusters 0, 1, 2, and 3 have medians around 200-250 and ranges from approximately 100 to 400.
- Violin Plot of Rolling7 by Cluster:** The y-axis is labeled 'Rolling7'. Cluster 4 has the highest median around 310 and a range from approximately 210 to 370. Cluster 3 has a median around 260 and a range from approximately 190 to 360. Clusters 0, 1, 2, and 5 have medians around 240-250 and ranges from approximately 200 to 300.
- Violin Plot of Lag1 by Cluster:** The y-axis is labeled 'Lag1'. Cluster 4 shows a very wide distribution with a median around 250 and a range from approximately -200 to 1200. Cluster 5 has a median around 200 and a range from approximately 100 to 300. Clusters 0, 1, 2, and 3 have medians around 250-300 and ranges from approximately 100 to 400.


```
C:\Users\Davon\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages
\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
    self._figure.tight_layout(*args, **kwargs)
C:\Users\Davon\AppData\Local\Temp\ipykernel_17252\3102834032.py:4: UserWarning: The figure layout has changed to tight
    plt.tight_layout()
```

Histogram of DailyExpense by Cluster

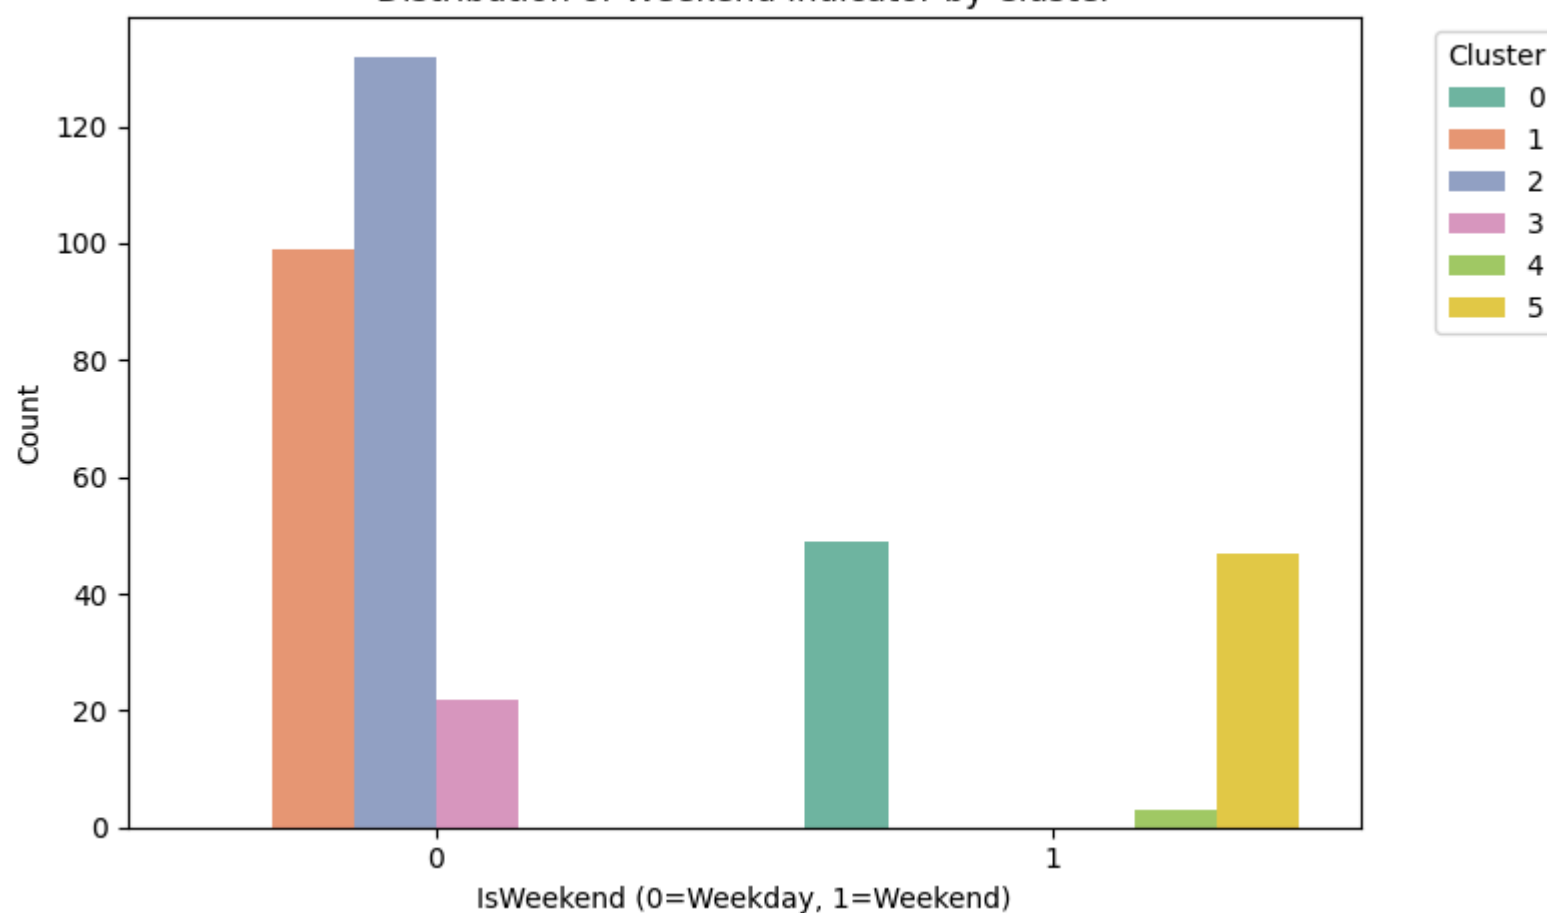


```
In [ ]: plt.figure(figsize=(12, 5))
sns.countplot(x='DayOfWeek', hue='GMM_Cluster', data=daily_expense, palette='Set2')
plt.title("Distribution of DayOfWeek by Cluster")
plt.xlabel("DayOfWeek (0=Monday, 6=Sunday)")
plt.ylabel("Count")
plt.legend(title="Cluster", bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 5))
sns.countplot(x='IsWeekend', hue='GMM_Cluster', data=daily_expense, palette='Set2')
plt.title("Distribution of Weekend Indicator by Cluster")
plt.xlabel("IsWeekend (0=Weekday, 1=Weekend)")
plt.ylabel("Count")
plt.legend(title="Cluster", bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```



Distribution of Weekend Indicator by Cluster



## Radar Chart for Cluster Mean Profiles

```
In [ ]: features_radar = ['DailyExpense', 'Rolling7', 'Lag1']
cluster_means = daily_expense.groupby('GMM_Cluster')[features_radar].mean().reset_index()

# Normalize each feature (min-max normalization) for fair comparison
def normalize(series):
    return (series - series.min()) / (series.max() - series.min())

for feat in features_radar:
    cluster_means[feat] = normalize(cluster_means[feat])

# Prepare the angles for the radar chart
num_vars = len(features_radar)
angles = np.linspace(0, 2 * np.pi, num_vars, endpoint=False).tolist()
angles += angles[:1]

# Plot radar chart
fig, ax = plt.subplots(figsize=(6, 6), subplot_kw=dict(polar=True))
for i, row in cluster_means.iterrows():
    values = row[features_radar].tolist()
    values += values[:1] # close the circle
    ax.plot(angles, values, label=f"Cluster {int(row['GMM_Cluster'])}")
    ax.fill(angles, values, alpha=0.25)
ax.set_thetagrids(np.degrees(angles[:-1]), features_radar)
ax.set_title("Radar Chart of Cluster Mean Profiles (Normalized)", fontsize=14)
ax.legend(loc='upper right', bbox_to_anchor=(1.3, 1.1))
plt.tight_layout()
plt.show()
```

Radar Chart of Cluster Mean Profiles (Normalized)

